

CSC 555 Assignment 2

Ronaldlee Ejalu

Reading: *Mining of Massive Datasets*: Chapter 2.

Hadoop: The Definitive Guide: Chapter 17 (Hive), Appendix A (file also available on D2L).

Supplemental document **UsingAmazonAWS.doc** and **Instructions_ReformatHDFS_Hive.doc**

The reformatting instructions are included in case you have to re-initialize your instance. You would **only** need it if you need to reformat your Hadoop set up.

Part 1

- 1) Describe how you would implement a MapReduce job consisting of Map and Reduce description. You can describe it in your own words or as pseudo-code. Keep in mind that map task reads the input file and produces (key, value) pairs. Reduce task takes a list of (key, value) pairs for each key and combines all values for each key.

Remember that Map operates on individual blocks and Reduce on individual keys with a set of values. Thus, for Mapper you need to state what your code does given a block of data and for Reduce you need to state what your reducer does for each key. You can assume that all of the columns accessed by the query exist in the original table.

- a)

```
SELECT First, SUM(Grade)
FROM Student
GROUP BY First;
```

Pseudo-code:

	Key	Value
Mapper	First	Grade
Reducer	First	SUM(Grade)

MapReduce works by breaking the processing into two phases: the map and the reduce phase as explained below:

The input to the mapper will be a row and it will produce key value pairs of First as the key and Grade as the value. Furthermore, First and Grade will be the key and value pairs respectively for each row that is processed by the mapper. The mapper emits the key value pairs (First: Grade), as shown above in the table, as its output.

The MapReduce framework receives output from the Mapper in form of keys and values and uses a deterministic hash function on the keys where it uses the value MOD N to distribute values to the different reducers. This process sorts and groups the key – value pairs by key, which is First. The list of grades will be the values for a particular

key, First. So, the key (First) and the values (list of grades) will act as an input to the reduce function or phase.

Using the list of grades per key, the reducer will generate the sum of grades with an output of key-value pairs corresponding to First and sum of Grades respectively.

- b) SELECT City, State, COUNT(DISTINCT Name)
FROM Student
GROUP BY City, State;
Pseudo-code:

	Key	Value
Mapper	City-State	DISTINCT NAME
Reducer	City-State	COUNT(DISTINCT NAME)

For each row the mapper processes, it emits an output of key value pairs.

The key will be a combination of City and State and the value will be the “Distinct” Name.

The output of the mapper i.e., City-State : {Distinct Name} will be received by the MapReduce framework, which uses a deterministic hash function on the keys by using the values of the hash function on keys MOD N to distribute values to the different reducers. This process sorts and groups the key – value pairs by the key, which is the combination of City and State. The reducers will receive the key, which is a combination of City and State and the values, which is a list of the “DISTINCT” Names as its inputs.

For each key, the reducer will count the distinct names in the list and the output City-State as the key and Count(DISTINCT NAME) as the value.

- 2) Suppose you are tasked with analysis of the company’s web server logs. The log dump contains a large amount of information with up to 7 different attributes (columns). You regularly run a Hadoop job to perform analysis pertaining to 3 specific attributes – TimeOfAccess, OriginOfAccess and FileName out of 7 total in the file.

- a) How would you attempt to speed up the regular execution of the query? (2-a is intentionally an open-ended question, there are several acceptable answers)

This will be achieved in a variety of ways:


- Rather than operating on individual rows of data at a time, the Hadoop job performs a batch of rows at a time. Operating on individual rows is like running a for loop on the individual elements of a big list which is slower.
- Out of the 7 attributes, we extract out the four attributes which are not needed for the job and only live TimeOfAccess, OriginOfAccess and FileName in the file on which we regularly run a Hadoop job to perform analysis. This will be done in the Map function.

- In the load function, I would implement the LoadPushDown Interface as a means of finding out which columns the query is asking for, this is where I would strictly specify the TimeOfAccess, OriginOfAccess and FileName attributes and leave out all the unnecessary attributes. This acts as an optimization technique for column-oriented storage, so that the loader only loads only those attributes needed by the query.
- b) If a Mapper task fails while processing a block of data – which node(s) where MapReduce framework will prefer to restart it?
MapReduce framework will restart the NameNode.
- c) If the job is executed with 3 Reducers
- i) How many files does the output generate?
The output generates 1 file.
 - ii) Suggest one possible hash function that may be used to assign keys to reducers.
The deterministic Hash function works on the key field of a record; use its value MOD N to distribute values over range 0 N-1. N is equal to the number of buckets, which correspond to reducers.
- 3) Consider a Hadoop job that processes an input data file of size equal to 65 disk blocks (65 different blocks, you can assume that HDFS replication factor is set to 1). The mapper in this job requires 1 minute to read and process a single block of data. For the purposes of this assignment, you can assume that the reduce part of this job takes zero time. You can also refer to the supplemental example on how to make this estimate.
- a) Approximately how long will it take to process the file if you only had one Hadoop worker node? You can assume that that only one mapper is created on every node.

Remember, the mapper in the job requires 1 minute to read and process a single block of data therefore for 65 blocks:

$$1 * 65 = 65 \text{ minutes.}$$

- b) 10 Hadoop worker nodes?



1	B1	B7	B13	B19	B25	B31	B37	B43	B49	B55
2	B2	B8	B14	B20	B26	B32	B38	B44	B50	B56
3	B3	B9	B15	B21	B27	B33	B39	B45	B51	B57
4	B4	B10	B16	B22	B28	B34	B40	B46	B52	B58
5	B5	B11	B17	B23	B29	B35	B41	B47	B53	B59
6	B6	B12	B18	B24	B30	B36	B42	B48	B54	B60
7	B61	B62	B63	B64	B65					

- So, 10 Hadoop worker nodes take 7 mins.

c) 30 Hadoop worker nodes?

	Node 1	Node 2	Node 3	Node 4	Node 5	Node 6	Node 7	Node 8	Node 9	Node10
1	B1	B3	B5	B7	B9	B11	B13	B15	B17	B19
2	B2	B4	B6	B8	B10	B12	B14	B16	B18	B20
3	B61	B62	B63	B64	B65					

Node 11	Node 12	Node 13	Node 14	Node 15	Node 16	Node 17	Node 18	Node 19	Node 20
B21	B23	B25	B27	B29	B31	B33	B35	B37	B39
B22	B24	B26	B28	B30	B32	B34	B36	B38	B40

Node 21	Node 22	Node 23	Node 24	Node 25	Node 26	Node 27	Node 28	Node 29	Node 30
B41	B43	B45	B47	B49	B51	B53	B55	B57	B59
B42	B44	B46	B48	B50	B52	B54	B56	B58	B60

- So, 30 Hadoop worker nodes are going to take 3 mins.

d) 100 Hadoop worker nodes?

- The 65 blocks will be distributed across the 65 worker nodes and 35 nodes will be left idle so it will take 1 min to process the 65 blocks on 100 Hadoop worker nodes.

e) Now suppose you were told that the replication factor has been changed to 3. That is, each block is stored in triplicate, but file size is still 65 blocks. Which of the answers (if any) in a)-d) above will have to change?

- Changing the replication factor to 3 has nothing to do with the time of the processing. A replication factor of 3 ensures a good balance among reliability, write bandwidth, read performance and block distribution across the cluster.

You can ignore the network transfer costs and other potential overheads as well as the possibility of node failure. State any assumptions you make.

- The communication cost, which often influences our choice of algorithm to use within a cluster computing environment.
- The importance of wall-clock time, the time it takes a parallel algorithm to finish; the operations should be divided fairly among tasks implying that the wall clock time would be approximately as small as it would be, given the number of compute nodes available.

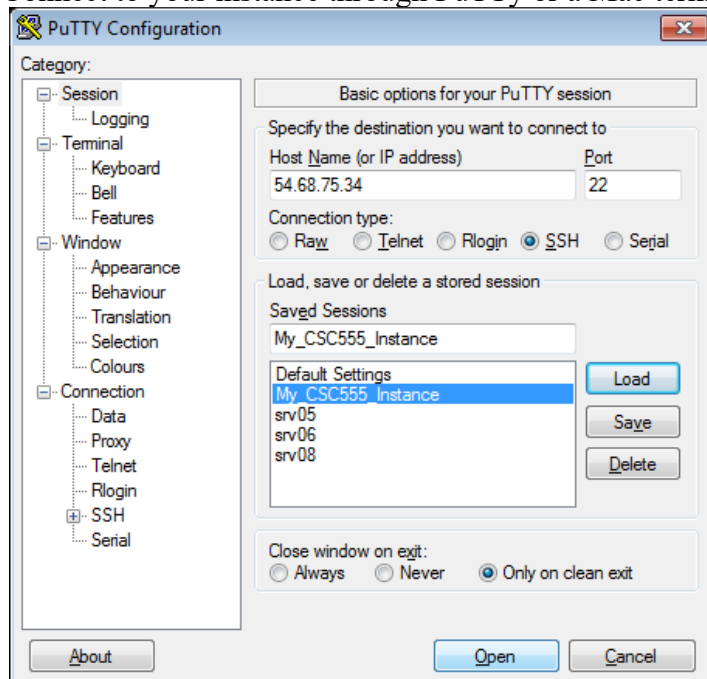
Part 2: Linux Intro

This part of the assignment will serve as an introduction to Linux. Make sure you go through the steps below and submit screenshots where requested – submit the entire screenshot of a command terminal.

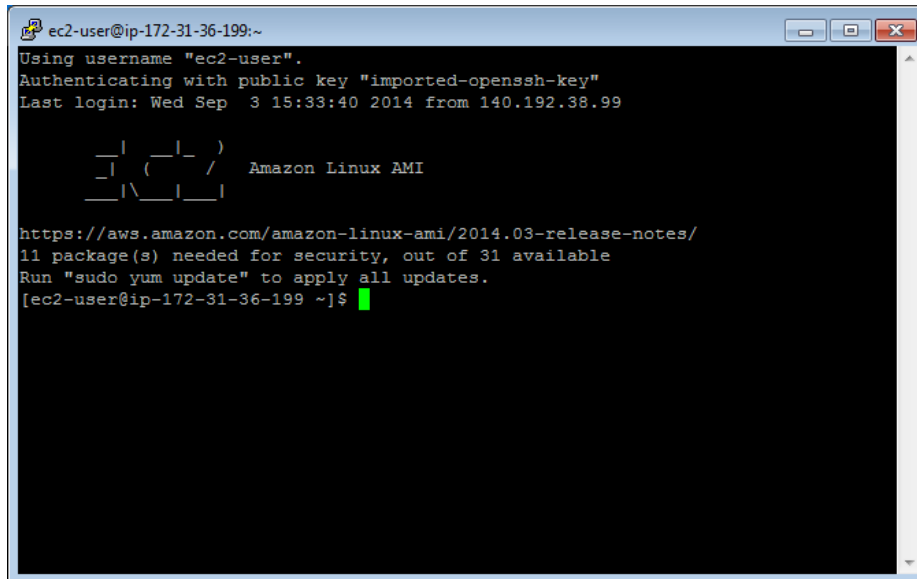
Use at least a t2.small instance or Hadoop may not run properly with insufficient memory. All Linux commands are in **Berlin Sans FB**. Do not type the “\$” symbol. The “\$” represents the prompt “[ec2-user@ip-xxx-xx-xx-xxx ~] \$ ” in your particular Linux instance.

0. Login to your Amazon EC2 Instance (NOTE: instructions on how to create a new instance and log in to it are provided in a separate file, UsingAmazonAWS.doc)

Connect to your instance through PuTTY or a Mac terminal.



On Windows, your instance would look similar to the following image. On a Mac, it would be an XTerm terminal with same text:

A screenshot of a Linux terminal window. The title bar shows 'ec2-user@ip-172-31-36-199:~'. The terminal text includes: 'Using username "ec2-user".', 'Authenticating with public key "imported-openssh-key"', 'Last login: Wed Sep 3 15:33:40 2014 from 140.192.38.99', an ASCII art logo for Amazon Linux AMI, a URL to AWS release notes, and a message about 11 security updates. The prompt is '[ec2-user@ip-172-31-36-199 ~]\$' with a green cursor.

```
ec2-user@ip-172-31-36-199:~  
Using username "ec2-user".  
Authenticating with public key "imported-openssh-key"  
Last login: Wed Sep 3 15:33:40 2014 from 140.192.38.99  
  
  _|  _|_ )  
  _| ( _| /  Amazon Linux AMI  
  _|\_|_|_|  
  
https://aws.amazon.com/amazon-linux-ami/2014.03-release-notes/  
11 package(s) needed for security, out of 31 available  
Run "sudo yum update" to apply all updates.  
[ec2-user@ip-172-31-36-199 ~]$
```

1. Create a text file.

Instructions for 3 different text editors are provided below. You only need to choose one editor that you prefer. **nano** is a more basic text editor, and is much easier to start. **vim** and **emacs** are more advanced and rely on keyboard shortcuts quite a bit and thus have a steeper learning curve.

- **Tip:** To paste into Linux terminal you can use right-click. To copy from the Linux terminal, you only need to highlight the text that you want to copy with your mouse. Also please remember that Linux is case-sensitive, which means **Nano** and **nano** are not equivalent.

Nano Instructions(Option 1):

```
$ nano myfile.txt
```

```
ec2-user@ip-172-31-36-199:~  
Using username "ec2-user".  
Authenticating with public key "imported-openssh-key"  
Last login: Wed Sep  3 15:33:40 2014 from 140.192.38.99  
  
  _ | _ | _ )  
  _ | ( _ /   Amazon Linux AMI  
  _ | \ _ | _ |  
  
https://aws.amazon.com/amazon-linux-ami/2014.03-release-notes/  
11 package(s) needed for security, out of 31 available  
Run "sudo yum update" to apply all updates.  
[ec2-user@ip-172-31-36-199 ~]$ nano myfile.txt
```

Type something into the file: “This is my text file for CSC555.”

```
ec2-user@ip-172-31-36-199:~  
GNU nano 2.0.9      File: myfile.txt  
  
This is my text file for CSC555.  
  
^G Get Help  ^O WriteOut  ^R Read File  ^Y Prev Page  ^K Cut Text   ^C Cur Pos  
^X Exit      ^J Justify   ^W Where Is   ^V Next Page  ^U UnCut Text ^T To Spell
```

Save changes: Ctrl-o and hit Enter.

Exit: Ctrl-x

Emacs Instructions (Option 2):

You will need to install emacs.

\$ sudo yum install emacs

Type “y” when asked if this is OK to install.

\$ emacs myfile.txt

Type something into the file: “This is my text file for CSC555.”

Save changes: Ctrl-x, Ctrl-s

Exit: Ctrl-x Ctrl-z

Vim Instructions(Option 3):

- NOTE: When **vim** opens, you are in *command* mode. Any key you enter will be bound to a command instead of inserted into the file. To enter *insert* mode press the key “i”. To save a file or exit you will need to hit Esc to get back into *command* mode.

\$ vim myfile.txt

Type “i” to enter *insert* mode

Type something into the file: “This is my text file for CSC555.”

Save changes: hit Esc to enter *command* mode then type “:w”

Exit: (still in *command* mode) type “:x”

Confirm your file has been saved by listing the files in the working directory.

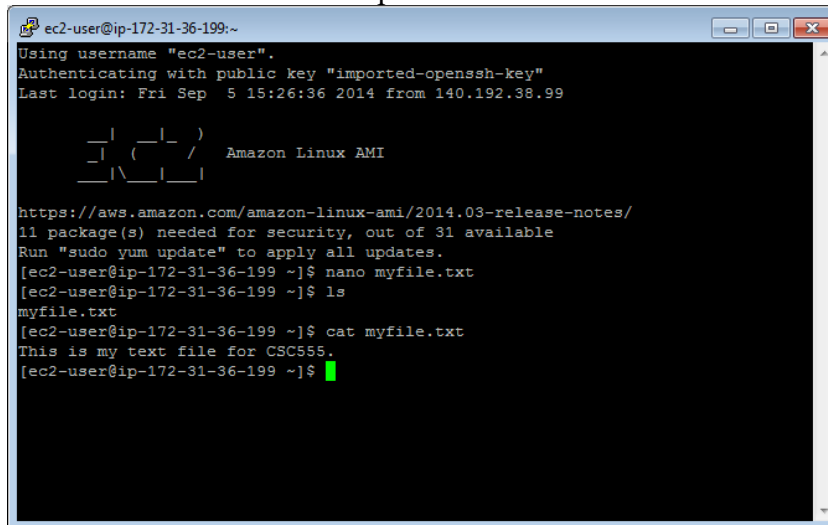
\$ ls

You should see your file.

Display the contents of the file on the screen.

\$ cat myfile.txt

Your file contents should be printed to the terminal.

A screenshot of a terminal window titled 'ec2-user@ip-172-31-36-199:~'. The terminal shows the following sequence of commands and outputs: 1. Login banner for Amazon Linux AMI. 2. A URL to AWS release notes. 3. A message about security updates. 4. Command '[ec2-user@ip-172-31-36-199 ~]\$ nano myfile.txt' followed by the file being created. 5. Command '[ec2-user@ip-172-31-36-199 ~]\$ ls' showing 'myfile.txt' in the directory. 6. Command '[ec2-user@ip-172-31-36-199 ~]\$ cat myfile.txt' showing the output 'This is my text file for CSC555.' 7. The prompt returns to '[ec2-user@ip-172-31-36-199 ~]\$'.

- **Tip:** Linux will fill in partially typed commands if you hit Tab.

\$cat myfi

Hit Tab and “myfi” should be completed to “myfile.txt”. If there are multiple completion options, hit Tab twice and a list of all possible completions will be printed. This also applies to commands themselves, i.e. you can type in **ca** and see all possible commands that begin with ca.

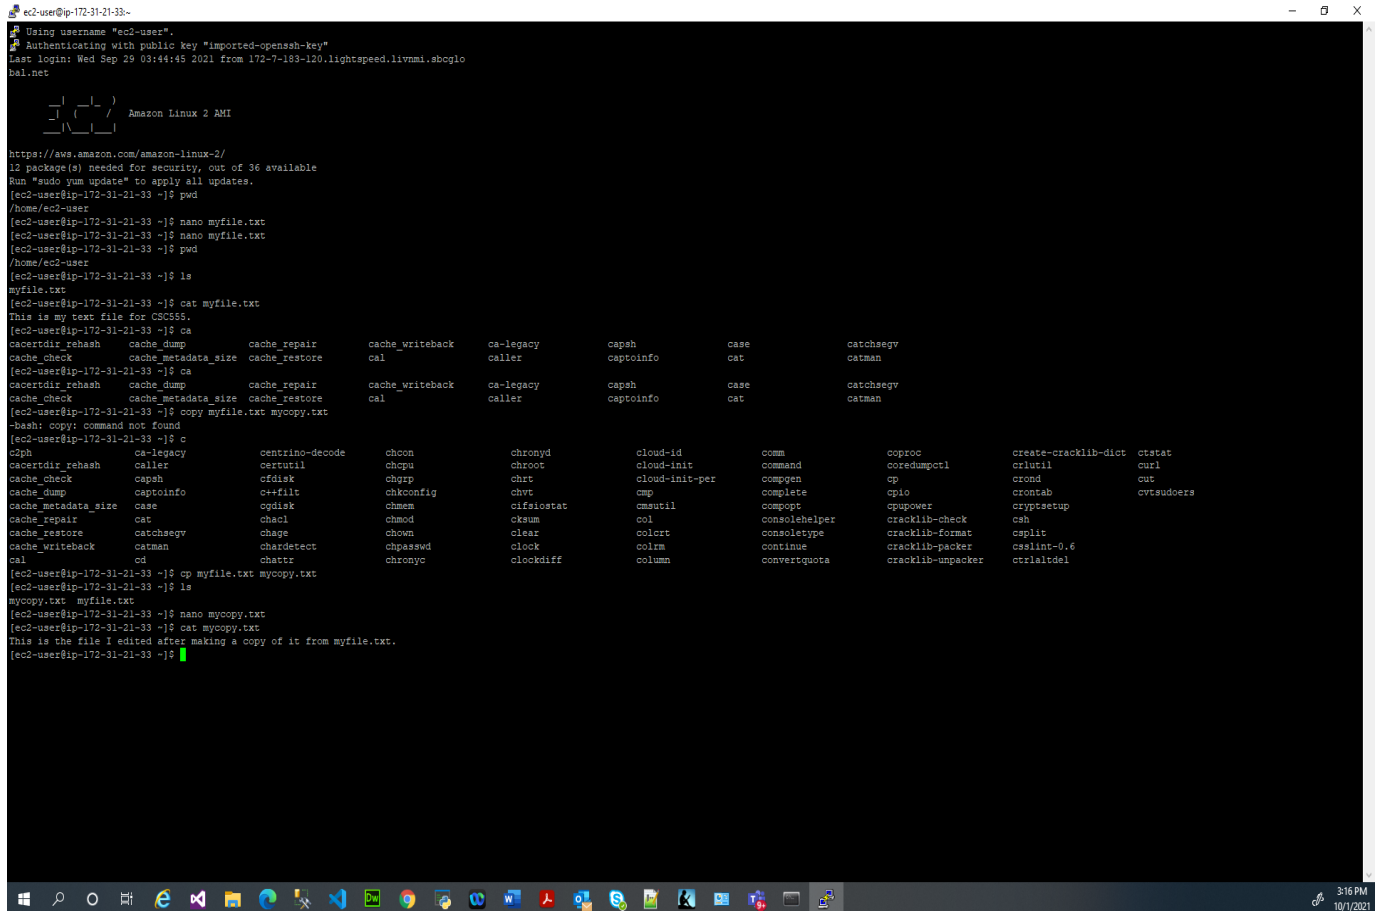
2. Copy your file.

Make a copy.

\$ cp myfile.txt mycopy.txt

Confirm this file has been created by listing the files in the working directory.
Edit this file so it contains different text than the original file using the text editor instructions,
and confirm your changes by displaying the contents of the file on the screen.

SUBMIT: Take a screen shot of the contents of your copied file displayed on the terminal screen.



```
ec2-user@ip-172-31-21-33:~$ Using username "ec2-user".
ec2-user@ip-172-31-21-33:~$ Authenticating with public key "imported-openssh-key"
ec2-user@ip-172-31-21-33:~$ Last login: Wed Sep 29 03:44:45 2021 from 172-7-183-120.lightspeed.livnmi.sbcoglo
ec2-user@ip-172-31-21-33:~$ bal.net

  _ _ _ _ _
 _ | ( _ | /
 _/ \_/_/_/_/   Amazon Linux 2 RMI

https://aws.amazon.com/amazon-linux-2/
12 package(s) needed for security, out of 36 available
Run "sudo yum update" to apply all updates.
ec2-user@ip-172-31-21-33 ~$ pwd
/home/ec2-user
ec2-user@ip-172-31-21-33 ~$ nano myfile.txt
ec2-user@ip-172-31-21-33 ~$ nano myfile.txt
ec2-user@ip-172-31-21-33 ~$ pwd
/home/ec2-user
ec2-user@ip-172-31-21-33 ~$ ls
myfile.txt
ec2-user@ip-172-31-21-33 ~$ cat myfile.txt
This is my text file for CSC555.
ec2-user@ip-172-31-21-33 ~$ ca
cacertdir_rehash  cache_dump  cache_repair  cache_writeback  ca-legacy  capah  case  catchsegv
cache_check  cache_metadata_size  cache_restore  cal  caller  captinfo  cat  catman
ec2-user@ip-172-31-21-33 ~$ ca
cacertdir_rehash  cache_dump  cache_repair  cache_writeback  ca-legacy  capah  case  catchsegv
cache_check  cache_metadata_size  cache_restore  cal  caller  captinfo  cat  catman
ec2-user@ip-172-31-21-33 ~$ cp myfile.txt mycopy.txt
-bash: copy: command not found
ec2-user@ip-172-31-21-33 ~$ c
c2ph  centrino-decode  chcon  chroot  cloud-id  comm  coppers  create-cracklib-dict  cstat
cacertdir_rehash  caller  certutil  chcpu  cloud-init  command  coredumpctl  curl
cache_check  capah  cfdisk  chgrp  cloud-init-per  compgen  cp  cron  cron  cron  cron
cache_dump  captinfo  c++filt  chkconfig  chvt  cmp  complete  cpio  cronab  cvsdoers
cache_metadata_size  case  cgdisk  chmem  cifsioatst  cmsutil  compopt  cpupower  cryptsetup
cache_repair  cat  chacl  chmod  cksun  col  consolehelper  cracklib-check  csh
cache_restore  catchsegv  chase  chown  clear  colcat  consolotype  cracklib-format  csplit
cache_writeback  catman  chardetct  chpasswd  clock  colrm  continue  cracklib-packer  csllint-0.6
cal  cd  chatrr  chronyc  clockdiff  column  convertquota  cracklib-umpacker  ctrlaltdel
ec2-user@ip-172-31-21-33 ~$ cp myfile.txt mycopy.txt
ec2-user@ip-172-31-21-33 ~$ ls
mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33 ~$ nano mycopy.txt
ec2-user@ip-172-31-21-33 ~$ cat mycopy.txt
This is the file I edited after making a copy of it from myfile.txt.
ec2-user@ip-172-31-21-33 ~$
```

3. Delete a file

Make a copy to delete.
\$ cp myfile.txt filetodelete.txt
\$ ls

Remove the file.
\$ rm filetodelete.txt
\$ ls

4. Create a directory to put your files.

Make a directory.
\$ mkdir CSC555

Change the current directory to your new directory.
\$cd CSC555

Print your current working directory
\$pwd

5. Move your files to your new directory.

Return to your home directory.

```
$cd
    OR
$cd ..
    OR
$ cd /home/ec2-user/
```

- NOTE: **cd** will always take you to your home directory. **cd ..** will move you up one directory level (to the parent). Your home directory is `"/home/[user name]"`, `/home/ec2-user` in our case

Move your files to your new directory.

```
$ mv myfile.txt CSC555/
$ mv mycopy.txt CSC555/
```

Change the current directory to CSC555 and list the files in this directory.

SUBMIT: Take a screen shot of the files listed in the CSC555 directory.

[illegible]

6. Zip and Unzip your files.

Zip the files.

```
$ zip myzipfile mycopy.txt myfile.txt
```

OR

```
$ zip myzipfile *
```

- NOTE: * is the wildcard symbol that matches everything in current directory. If there should be any additional files in the current directory, they will also be placed into the zip archive. Wildcard can also be used to match files selectively. For example **zip myzipfile my*** will zip-up all files beginning with “my” in the current directory.

Move your zip file to your home directory.

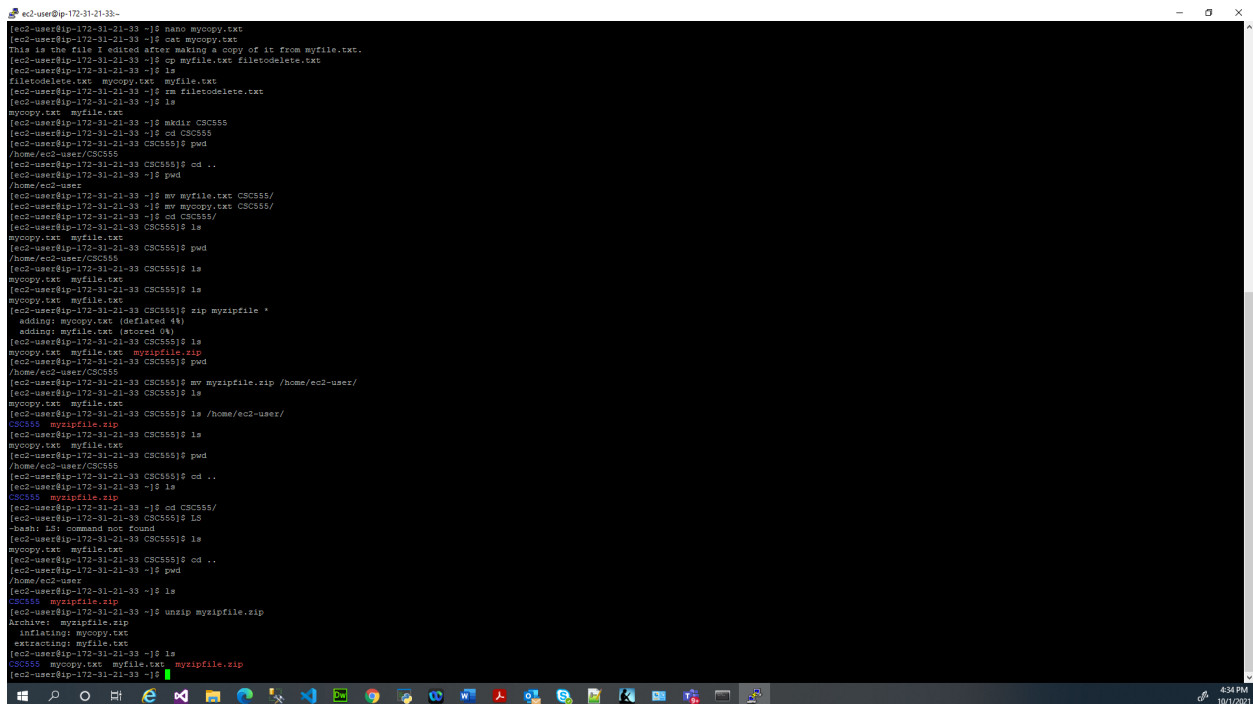
```
$ mv myzipfile.zip /home/ec2-user/
```

Return to your home directory.

Extract the files.

```
$ unzip myzipfile.zip
```

SUBMIT: Take a screen shot of the screen after this command.



```
ec2-user@ip-172-31-21-33:~$ nano mycopy.txt
ec2-user@ip-172-31-21-33:~$ cat mycopy.txt
This is the file I edited after making a copy of it from myfile.txt.
ec2-user@ip-172-31-21-33:~$ cp myfile.txt filedelete.txt
ec2-user@ip-172-31-21-33:~$ ls
filedelete.txt  mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33:~$ rm filedelete.txt
ec2-user@ip-172-31-21-33:~$ ls
mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33:~$ mkdir CSC555
ec2-user@ip-172-31-21-33:~$ cd CSC555
ec2-user@ip-172-31-21-33:~/CSC555$ pwd
/home/ec2-user/CSC555
ec2-user@ip-172-31-21-33:~/CSC555$ cd ..
ec2-user@ip-172-31-21-33:~$ pwd
/home/ec2-user
ec2-user@ip-172-31-21-33:~$ mv myfile.txt CSC555/
ec2-user@ip-172-31-21-33:~$ mv mycopy.txt CSC555/
ec2-user@ip-172-31-21-33:~$ cd CSC555/
ec2-user@ip-172-31-21-33:~/CSC555$ ls
mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33:~/CSC555$ pwd
/home/ec2-user/CSC555
ec2-user@ip-172-31-21-33:~/CSC555$ ls
mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33:~/CSC555$ zip myzipfile *
adding: mycopy.txt (deflated 0%)
adding: myfile.txt (stored 0%)
ec2-user@ip-172-31-21-33:~/CSC555$ ls
mycopy.txt  myfile.txt  myzipfile.zip
ec2-user@ip-172-31-21-33:~/CSC555$ pwd
/home/ec2-user/CSC555
ec2-user@ip-172-31-21-33:~/CSC555$ mv myzipfile.zip /home/ec2-user/
ec2-user@ip-172-31-21-33:~/CSC555$ ls
mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33:~/CSC555$ ls /home/ec2-user/
CSC555  myzipfile.zip
ec2-user@ip-172-31-21-33:~/CSC555$ ls
mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33:~/CSC555$ pwd
/home/ec2-user/CSC555
ec2-user@ip-172-31-21-33:~/CSC555$ cd ..
ec2-user@ip-172-31-21-33:~$ ls
CSC555  myzipfile.zip
ec2-user@ip-172-31-21-33:~$ cd CSC555/
ec2-user@ip-172-31-21-33:~/CSC555$ ls
-bash: ls: command not found
ec2-user@ip-172-31-21-33:~/CSC555$ ls
mycopy.txt  myfile.txt
ec2-user@ip-172-31-21-33:~/CSC555$ cd ..
ec2-user@ip-172-31-21-33:~$ pwd
/home/ec2-user
ec2-user@ip-172-31-21-33:~$ ls
CSC555  myzipfile.zip
ec2-user@ip-172-31-21-33:~$ unzip myzipfile.zip
Archive:  myzipfile.zip
  inflating: mycopy.txt
  extracting: myfile.txt
ec2-user@ip-172-31-21-33:~$ ls
CSC555  mycopy.txt  myfile.txt  myzipfile.zip
ec2-user@ip-172-31-21-33:~$
```

7. Remove your CSC555 directory.

- **Warning:** Executing “rm -rf” has the potential to delete ALL files in a given directory, including sub-directories (“r” stands for recursive). You should use this command very carefully.

Delete your CSC555 directory.

```
$ rm -rf CSC555/
```

8. Download a file from the web.

Download the script for Monty Python and the Holy Grail.

```
$ wget http://www.textfiles.com/media/SCRIPTS/grail
```

The file should be saved as “grail” by default.

9. ls formats

List all contents of the current directory in long list format.

- **Note:** the option following “ls” is the character “l”; not “one”.

\$ ls -l

The 1st column gives information regarding file permissions (which we will discuss in more detail later). For now, note that the first character of the 10 total will be “-“ for normal files and “d” for directories. The 2nd Column is the number of links to the file. The 3rd and 4th columns are the owner and the group of the file. The 5th column displays the size of the file in bytes. The 6th column is the date and time the file was last modified. The 7th column is the file or directory name.

List all contents of the current directory in long list and human readable formats. “-h” will put large files in more readable units than bytes.

\$ ls -lh

SUBMIT: The size of the grail file.

The size of the grail file is 73K as shown below:

```

root@kali:~# cd /tmp/172.31.21.33
/home/ec2-user/CSC355
[ec2-user@ip-172-31-21-33 CSC355]$ mv myfile.zip /home/ec2-user/
[ec2-user@ip-172-31-21-33 CSC355]$ ls
mycopy.txt  myfile.txt
[ec2-user@ip-172-31-21-33 CSC355]$ ls /home/ec2-user/
CSC355  myfile.zip
[ec2-user@ip-172-31-21-33 CSC355]$ ls
mycopy.txt  myfile.txt
[ec2-user@ip-172-31-21-33 CSC355]$ pwd
/home/ec2-user/CSC355
[ec2-user@ip-172-31-21-33 CSC355]$ cd ..
[ec2-user@ip-172-31-21-33 ~]$ ls
CSC355  myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ cd CSC355/
[ec2-user@ip-172-31-21-33 CSC355]$ ls
bash: ls: command not found
[ec2-user@ip-172-31-21-33 CSC355]$ ls
mycopy.txt  myfile.txt
[ec2-user@ip-172-31-21-33 CSC355]$ cd ..
[ec2-user@ip-172-31-21-33 ~]$ pwd
/home/ec2-user
[ec2-user@ip-172-31-21-33 ~]$ ls
CSC355  myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ unzip myfile.zip
Archive:  myfile.zip
  inflating: mycopy.txt
    extracting: myfile.txt
[ec2-user@ip-172-31-21-33 ~]$ ls
CSC355  mycopy.txt  myfile.txt  myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ ls
CSC355  mycopy.txt  myfile.txt  myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ cd /usr/CSC355/
[ec2-user@ip-172-31-21-33 ~]$ ls
bash: ls: command not found
[ec2-user@ip-172-31-21-33 ~]$ ls
mycopy.txt  myfile.txt  myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ wget http://www.textfiles.com/media/SCRIPTS/grail
--2021-10-01 20:41:13--  http://www.textfiles.com/media/SCRIPTS/grail
Resolving www.textfiles.com (www.textfiles.com)... 208.86.224.90
Connecting to www.textfiles.com (www.textfiles.com)[208.86.224.90]:80... connected.
HTTP request sent, awaiting response... 200 OK
length: 74635 (73K)
saving to 'grail'

100K[----->>>] 74,635  --K/s  in 0.05s

2021-10-01 20:41:13 (1.36 MB/s) - 'grail' saved [74635/74635]

[ec2-user@ip-172-31-21-33 ~]$ ls
grail  mycopy.txt  myfile.txt  myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ ls -l
total 56
-rw-rw-r-- 1 ec2-user ec2-user 74635 Aug 9 20:00 grail
-rw-rw-r-- 1 ec2-user ec2-user 60 Oct 1 15:15 mycopy.txt
-rw-rw-r-- 1 ec2-user ec2-user 33 Oct 1 18:25 myfile.txt
-rw-rw-r-- 1 ec2-user ec2-user 417 Oct 1 20:28 myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ ls -lh
total 88K
-rw-rw-r-- 1 ec2-user ec2-user 73K Aug 9 20:00 grail
-rw-rw-r-- 1 ec2-user ec2-user 69 Oct 1 15:15 mycopy.txt
-rw-rw-r-- 1 ec2-user ec2-user 33 Oct 1 18:25 myfile.txt
-rw-rw-r-- 1 ec2-user ec2-user 417 Oct 1 20:28 myfile.zip
[ec2-user@ip-172-31-21-33 ~]$ ls

```

10. More on viewing files.

If you issue “cat grail”, the contents of grail will be printed. However, this file is too large to fit on the screen.

Show the grail file one page at a time.

\$ more grail

Hit the spacebar to go to the next page. Type “b” to go page up, hit “space” key to go page down. Type “q” to quit.

OR

\$ less grail

Less has more options than *more* (“*less* is more and *more* is less”). You can now use the keyboard Arrows and Page Up/Down to scroll. You can type “h” for help, which will display additional options.

View the line numbers in the grail file. The *cat* command has the -n option, which prints line numbers, but you may also want to use *more* to view the file one page at a time. A solution is to *pipe* the output from *cat* to *more*. A *pipe* redirects the output from one program to another program for further processing, and it is represented with “|”.

\$ cat -n grail | more

Redirect the standard output (stdout) of a command.

\$ cat myfile.txt > redirect1.txt

\$ ls -lh > redirect2.txt

Append the stdout to a file.

\$ cat mycopy.txt >> myfile.txt

mycopy.txt will be appended to myfile.txt.

- **Note:** “cat mycopy.txt > myfile.txt” will overwrite myfile.txt with the contents output by “cat mycopy.txt”. Thus using >> is crucial if you want to preserve the existing file contents.

11. Change access permissions to objects with the *change mode* command.

The following represent roles:

u – user, g – group, o – others, a - all

The following represent permissions:

r – read, w – write, x – execute

Remove the read permission for your user on a file.

\$ chmod u-r myfile.txt

Try to read this file. You should receive a “permission denied” message because you are the user who owns the file.

SUBMIT: The screenshot of the permission denied error

```
ec2-user@ip-172-31-21-33:~$ cat myfile.txt
2042 REVERSE: With whom?
2043 INSPECTOR: Which one?
2044 OFFICER #1: Oh-- this one.
2045 INSPECTOR: Come on. Put him in the van.
2046 OFFICER #2: Get a blanket.
2047 OFFICER #1: We have no hospital.
2048 RANDOM: Ahh.
2049 [squeak]
2050 RANDOM: Oh.
2051 OFFICER #1: Come on. Back. Riight back. Come on!
2052 OFFICER #2: Run along! Run along!
2053 OFFICER #1: Pull that off. My, that's an offensive weapon, that is.
2054 OFFICER #2: Come on. Back with 'em. Back. Right. Come along.
2055 INSPECTOR: Everything?
2056 [squeak]
2057 OFFICER #1: All right, sonny. That's enough. Just pack that in.
2058 [crash]
2059 CAMERAMAN: Christ!
2060
2061
2062
2063
2064
THE END
ec2-user@ip-172-31-21-33 ~$ ls
grail  mycopy.txt  myfile.txt  myzipfile.zip
ec2-user@ip-172-31-21-33 ~$ cat myfile.txt > redirect1.txt
ec2-user@ip-172-31-21-33 ~$ ls
grail  mycopy.txt  myfile.txt  myzipfile.zip  redirect1.txt
ec2-user@ip-172-31-21-33 ~$ cat redirect1.txt
This is my text file for CSC555.
ec2-user@ip-172-31-21-33 ~$ ls -lh > redirect2.txt
ec2-user@ip-172-31-21-33 ~$ cat redirect2.txt
total 92K
-rw-rw-r-- 1 ec2-user ec2-user 73K Aug 9 2000 grail
-rw-rw-r-- 1 ec2-user ec2-user 49 Oct 1 19:15 mycopy.txt
-rw-rw-r-- 1 ec2-user ec2-user 33 Oct 1 18:25 myfile.txt
-rw-rw-r-- 1 ec2-user ec2-user 417 Oct 1 20:28 myzipfile.zip
-rw-rw-r-- 1 ec2-user ec2-user 33 Oct 1 21:06 redirect1.txt
-rw-rw-r-- 1 ec2-user ec2-user 0 Oct 1 21:06 redirect2.txt
ec2-user@ip-172-31-21-33 ~$ cat mycopy.txt > myfile.txt
ec2-user@ip-172-31-21-33 ~$ cat myfile.txt
This is the file I edited after making a copy of it from myfile.txt.
ec2-user@ip-172-31-21-33 ~$ ls -l
total 96
-rw-rw-r-- 1 ec2-user ec2-user 74635 Aug 9 2000 grail
-rw-rw-r-- 1 ec2-user ec2-user 49 Oct 1 19:15 mycopy.txt
-rw-rw-r-- 1 ec2-user ec2-user 102 Oct 1 21:50 myfile.txt
-rw-rw-r-- 1 ec2-user ec2-user 417 Oct 1 20:28 myzipfile.zip
-rw-rw-r-- 1 ec2-user ec2-user 33 Oct 1 21:05 redirect1.txt
-rw-rw-r-- 1 ec2-user ec2-user 368 Oct 1 21:06 redirect2.txt
ec2-user@ip-172-31-21-33 ~$ chmod u-r myfile.txt
ec2-user@ip-172-31-21-33 ~$ ls -l
total 96
-rw-rw-r-- 1 ec2-user ec2-user 74635 Aug 9 2000 grail
-rw-rw-r-- 1 ec2-user ec2-user 49 Oct 1 19:15 mycopy.txt
-rw-rw-r-- 1 ec2-user ec2-user 102 Oct 1 21:50 myfile.txt
-rw-rw-r-- 1 ec2-user ec2-user 417 Oct 1 20:28 myzipfile.zip
-rw-rw-r-- 1 ec2-user ec2-user 33 Oct 1 21:05 redirect1.txt
-rw-rw-r-- 1 ec2-user ec2-user 368 Oct 1 21:06 redirect2.txt
ec2-user@ip-172-31-21-33 ~$ cat myfile.txt
cat: myfile.txt: Permission denied
ec2-user@ip-172-31-21-33 ~$
```

Give your user read permission on a file. Use the same file you removed the read permission from.

\$ chmod u+r myfile.txt

You should now be able to read this file again.

12. Python examples

Install Python if it is not available on your machine.

\$ sudo yum install python

Create a Python file. These instructions will use Emacs as a text editor, but you can still chose the text editor you want.

\$ emacs lucky.py

(Write a simple Python program)

```
print "*" * 25
print "My Lucky Numbers".rjust(20)
print "*" * 25

for i in range(10):
    lucky_nbr = (i + 1) * 2
    print "My lucky number is %s!" % lucky_nbr
```

Run your Python program.

\$ python lucky.py

Redirect your output to a file
\$ python lucky.py > lucky.txt

Pipe the stdout from lucky.py to another Python program that will replace “is” with “was”.
\$ emacs was.py

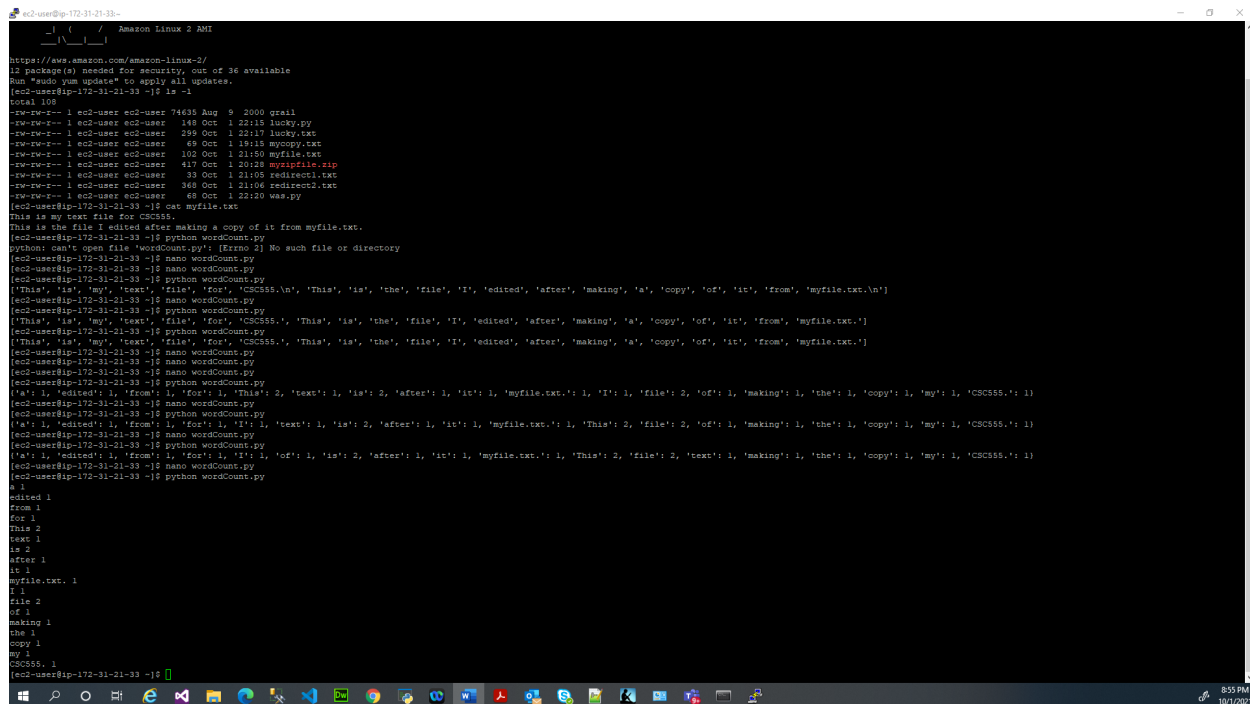
```
import sys
```

```
for line in sys.stdin:  
    print line.replace("is", "was")
```

\$ python lucky.py | python was.py

Write python code to read a text file (you can use myfile.txt) and output a word count total for each word with the number of times that word occurs in the entire file. That is, if the file has the word “Hadoop” occurs in the file 5 times, your code should print “Hadoop 5”. It should output the count of all words occurring in a file.

SUBMIT: The screen output from running your python code and a copy of your python code. Homework submissions without code will receive no credit.



```
ec2-user@ip-172-31-21-33:~$ cat myfile.txt
This is my text file for CSC355.
This is the file I edited after making a copy of it from myfile.txt.
ec2-user@ip-172-31-21-33:~$ python wordCount.py
python: can't open file 'wordCount.py': [Errno 2] No such file or directory
ec2-user@ip-172-31-21-33:~$ nano wordCount.py
ec2-user@ip-172-31-21-33:~$ python wordCount.py
['This', 'is', 'my', 'text', 'file', 'for', 'CSC355.', 'This', 'is', 'the', 'file', 'I', 'edited', 'after', 'making', 'a', 'copy', 'of', 'it', 'from', 'myfile.txt.\n']
ec2-user@ip-172-31-21-33:~$ nano wordCount.py
ec2-user@ip-172-31-21-33:~$ python wordCount.py
['This', 'is', 'my', 'text', 'file', 'for', 'CSC355.', 'This', 'is', 'the', 'file', 'I', 'edited', 'after', 'making', 'a', 'copy', 'of', 'it', 'from', 'myfile.txt.\n']
ec2-user@ip-172-31-21-33:~$ nano wordCount.py
ec2-user@ip-172-31-21-33:~$ python wordCount.py
['This', 'is', 'my', 'text', 'file', 'for', 'CSC355.', 'This', 'is', 'the', 'file', 'I', 'edited', 'after', 'making', 'a', 'copy', 'of', 'it', 'from', 'myfile.txt.\n']
ec2-user@ip-172-31-21-33:~$ nano wordCount.py
ec2-user@ip-172-31-21-33:~$ python wordCount.py
('a': 1, 'edited': 1, 'from': 1, 'for': 1, 'This': 2, 'text': 1, 'is': 2, 'after': 1, 'it': 1, 'myfile.txt': 1, 'I': 1, 'file': 2, 'of': 1, 'making': 1, 'the': 1, 'copy': 1, 'my': 1, 'CSC355.': 1)
ec2-user@ip-172-31-21-33:~$ nano wordCount.py
ec2-user@ip-172-31-21-33:~$ python wordCount.py
('a': 1, 'edited': 1, 'from': 1, 'for': 1, 'I': 1, 'text': 1, 'is': 2, 'after': 1, 'it': 1, 'myfile.txt': 1, 'This': 2, 'file': 2, 'of': 1, 'making': 1, 'the': 1, 'copy': 1, 'my': 1, 'CSC355.': 1)
ec2-user@ip-172-31-21-33:~$ nano wordCount.py
ec2-user@ip-172-31-21-33:~$ python wordCount.py
('a': 1, 'edited': 1, 'from': 1, 'for': 1, 'I': 1, 'of': 1, 'is': 2, 'after': 1, 'it': 1, 'myfile.txt': 1, 'This': 2, 'file': 2, 'text': 1, 'making': 1, 'the': 1, 'copy': 1, 'my': 1, 'CSC355.': 1)
ec2-user@ip-172-31-21-33:~$ nano wordCount.py
ec2-user@ip-172-31-21-33:~$ python wordCount.py
a 1
edited 1
from 1
for 1
This 2
text 1
is 2
after 1
it 1
myfile.txt. 1
I 1
file 2
of 1
making 1
the 1
copy 1
my 1
my 1
CSC355. 1
ec2-user@ip-172-31-21-33:~$
```

A copy of my python code:

```
def readFile(fileName):  
    """Function that returns a list of words after reading a file"""  
    bagOfWordsL = []  
    with open(fileName) as f:
```

```

        # read the file and generate a list of strings
        lines = f.readlines()

        for i in range(len(lines)):
            # a list of words is created by splitting the
string            splittedWordsL = lines[i].split(' ')

            # loop through the list of words and add words to
the bag of words list            for splittedItem in range(len(splittedWordsL)):
                # if there are any empty spaces just pass
                if len(splittedWordsL[splittedItem]) == 0:
                    pass
                else:
                    # used a replace function to
remove all line breaks.

bagOfWordsL.append(splittedWordsL[splittedItem].replace('\n',''))

        return bagOfWordsL            # return a list of bag of words.

def wordCnt(listOfWords):
    """Function that generates a dictionary of words with their frequency
of occurrence"""
    wordCnt = {}
    # define an empty dictionary
    for i in range(len(listOfWords)):
    # iterate through the list of words and count the number of words
        if listOfWords[i] in wordCnt.keys():
            wordCnt[listOfWords[i]] = wordCnt[listOfWords[i]] + 1
        else:
            wordCnt[listOfWords[i]] = 1
    return wordCnt
    # return a dictionary of the different words with their frequencies.

listOfBagOfWords = readFile('myfile.txt')
# return a list of words
wordsCntDict = wordCnt(listOfBagOfWords)
# uses a helper function that returns a dictionary of the different
words with their frequencies

# sort the dictionary
sortedWordCntDict = {key: value for (key, value) in
sorted(wordsCntDict.items(), key=lambda wordCnt: wordCnt[1], reverse =
True)}
for key in sortedWordCntDict:
    print('%s %s' %(key, sortedWordCntDict[key]))

```

Part 3: Wordcount

For this part of the assignment, you will run wordcount on a single-node Hadoop instance. I am going to provide detailed instructions to help you get Hadoop running. The instructions are

following Hadoop: The Definitive Guide instructions presented in Appendix A: Installing Apache Hadoop.

You can download 2.6.4 from here. You can copy-paste these commands (right-click in PuTTY to paste, but please watch out for error messages and run commands one by one)

Install ant to list java processes
sudo yum install ant

(wget command stands for “web get” and lets you download files to your instance from a URL link)

wget [http:// dbgroup.cdm.depaul.edu/Courses/CSC555/hadoop-2.6.4.tar.gz](http://dbgroup.cdm.depaul.edu/Courses/CSC555/hadoop-2.6.4.tar.gz)

(unpack the archive)
tar xzf hadoop-2.6.4.tar.gz

Modify the conf/hadoop-env.sh to add to it the JAVA_HOME configuration
You can open it by running (using nano or your favorite editor instead of nano).

nano hadoop-2.6.4/etc/hadoop/hadoop-env.sh

Note that the # comments out the line, so you would comment out the original JAVA_HOME line replacing it by the new one as below.

NOTE: you would need to determine the correct Java configuration line by executing the following (underlined) command

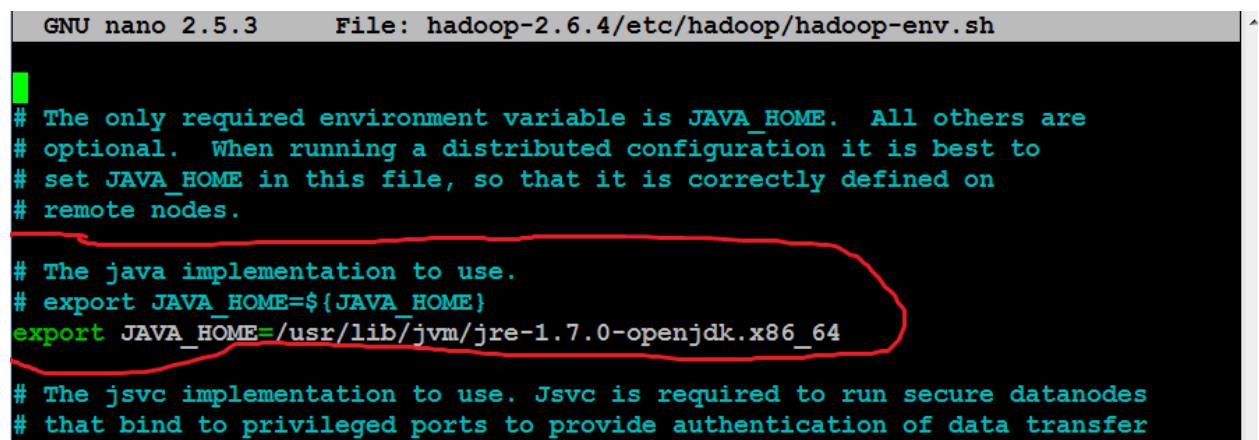
[ec2-user@ip-172-31-16-63 ~]\$ readlink -f \$(which java)

which will output something like:

/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.191.b12-0.amzn2.x86_64/jre/bin/java

In my case, Java home is at (remove the bin/java from the path):

/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.191.b12-0.amzn2.x86_64/jre/



```
GNU nano 2.5.3      File: hadoop-2.6.4/etc/hadoop/hadoop-env.sh

# The only required environment variable is JAVA_HOME.  All others are
# optional.  When running a distributed configuration it is best to
# set JAVA_HOME in this file, so that it is correctly defined on
# remote nodes.

# The java implementation to use.
# export JAVA_HOME=${JAVA_HOME}
export JAVA_HOME=/usr/lib/jvm/jre-1.7.0-openjdk.x86_64

# The jsvc implementation to use. Jsvc is required to run secure datanodes
# that bind to privileged ports to provide authentication of data transfer
```

modify the .bashrc file to add these two lines:

export HADOOP_HOME=~/.hadoop-2.6.4

export PATH=\$PATH:\$HADOOP_HOME/bin:\$HADOOP_HOME/sbin

.bashrc file contains environment settings to be configured automatically on each login. You can open the .bashrc file by running
`nano ~/.bashrc`

```
# User specific aliases and functions

export HADOOP_HOME=~/.hadoop-2.6.4
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin
```

To immediately refresh the settings (that will be automatic on next login), run
`source ~/.bashrc`

Next, follow the instructions for Pseudodistributed Mode for all 4 files.

(to edit the first config file)
`nano hadoop-2.6.4/etc/hadoop/core-site.xml`

Make sure you paste the settings between the <configuration> and </configuration> tags, like in the screenshot below. NOTE: The screenshot below is only one of the 4 files, all files are different. The contents of each file are described in the **Appendix A** in the Hadoop book, the relevant appendix is also included with the homework assignment. I am also including a .txt file (HadoopConfigurationText) so that it is easier to copy-paste.

```
!-- Put site-specific property overrides in this file. -->

<configuration>

  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost/</value>
  </property>

</configuration>
```

`nano hadoop-2.6.4/etc/hadoop/hdfs-site.xml`
(mapred-site.xml file is not there, run the following **single line** command to create it by copying from template. Then you can edit it as other files.)

`cp hadoop-2.6.4/etc/hadoop/mapred-site.xml.template hadoop-2.6.4/etc/hadoop/mapred-site.xml`

`nano hadoop-2.6.4/etc/hadoop/mapred-site.xml`
`nano hadoop-2.6.4/etc/hadoop/yarn-site.xml`

To enable passwordless ssh access (we will discuss SSH and public/private keys in class), run these commands:

`ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa`
`cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys`

test by running (and confirming yes to a one-time warning)
ssh localhost
exit

Format HDFS (i.e., first time initialize)

hdfs namenode -format

Start HDFS, Hadoop and history server (answer a 1-time yes if you asked about host authenticity)

start-dfs.sh
start-yarn.sh
mr-jobhistory-daemon.sh start historyserver

Verify if everything is running:
jps

(NameNode and DataNode are responsible for HDFS management; NodeManager and ResourceManager are serving the function similar to JobTracker and TaskTracker.)

Create a destination directory
hadoop fs -mkdir /data

Download a large text file using
wget <http://dbgroup.cdm.depaul.edu/Courses/CSC555/bioproject.xml>

Copy the file to HDFS for processing

hadoop fs -put bioproject.xml /data/

(you can optimally verify that the file was uploaded to HDFS by **hadoop fs -ls /data**)
Submit a screenshot of this command

```
ec2-user@ip-172-31-21-33:~$
drwxr-xr-x 3 ec2-user ec2-user 4096 Oct 2 03:10 logs
-rw-r--r-- 1 ec2-user ec2-user 103 Feb 15 2016 NGIPE.txt
-rw-r--r-- 1 ec2-user ec2-user 1366 Feb 12 2016 README.txt
drwxr-xr-x 2 ec2-user ec2-user 4096 Feb 12 2016 sbin
drwxr-xr-x 4 ec2-user ec2-user 31 Feb 12 2016 share
[ec2-user@ip-172-31-21-33 ~]$ ls -l /
total 29
lrwxrwxrwx 1 root root 7 Aug 24 19:05 bin -> usr/bin
drwxr-xr-x 4 root root 4096 Aug 24 19:06 boot
drwxr-xr-x 15 root root 2820 Oct 1 17:44 dev
drwxr-xr-x 88 root root 9192 Oct 2 02:18 etc
drwxr-xr-x 3 root root 22 Sep 29 03:15 home
lrwxrwxrwx 1 root root 7 Aug 24 19:05 lib -> usr/lib
lrwxrwxrwx 1 root root 9 Aug 24 19:05 lib64 -> usr/lib64
drwxr-xr-x 2 root root 6 Aug 24 19:04 local
drwxr-xr-x 2 root root 6 Apr 9 2015 media
drwxr-xr-x 2 root root 6 Apr 9 2015 mnt
drwxr-xr-x 4 root root 27 Aug 24 19:06 opt
drwxr-xr-x 102 root root 0 Oct 1 17:44 proc
drwxr-xr-x 3 root root 103 Sep 29 03:15 root
drwxr-xr-x 28 root root 960 Oct 2 02:18 run
lrwxrwxrwx 1 root root 9 Aug 24 19:05/sbin -> usr/sbin
drwxr-xr-x 2 root root 6 Apr 9 2015 srv
drwxr-xr-x 13 root root 0 Oct 1 17:44 sys
drwxr-xr-x 17 root root 4096 Oct 2 03:19 tmp
drwxr-xr-x 13 root root 155 Aug 24 19:05 var
drwxr-xr-x 49 root root 269 Sep 29 03:15 vvar
[ec2-user@ip-172-31-21-33 ~]$ ls -l /data
ls: cannot access /data: No such file or directory
[ec2-user@ip-172-31-21-33 ~]$ hadoop fs -mkdir /data
mkdir: '/data': File exists
[ec2-user@ip-172-31-21-33 ~]$ wget http://dbgroup.cdm.depaul.edu/Courses/CSC555/bioproject.xml
--2021-10-02 03:41:11-- http://dbgroup.cdm.depaul.edu/Courses/CSC555/bioproject.xml
Resolving dbgroup.cdm.depaul.edu (dbgroup.cdm.depaul.edu)... 140.192.248.250
Connecting to dbgroup.cdm.depaul.edu (dbgroup.cdm.depaul.edu)|140.192.248.250|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 23114903 (22.0M) [application/xml]
Saving to: 'bioproject.xml'

100%[=====] 231,149,003 73.0MB/s in 3.0s

2021-10-02 03:41:14 (73.0 MB/s) - 'bioproject.xml' saved [231149003/231149003]

[ec2-user@ip-172-31-21-33 ~]$ ls -l
total 41950
-rw-rw-r-- 1 ec2-user ec2-user 231149003 Sep 27 23:08 bioproject.xml
-rw-rw-r-- 1 ec2-user ec2-user 74635 Aug 9 2000 gra1l
drwxr-xr-x 10 ec2-user ec2-user 403 Oct 2 03:29 hadoop-2.6.4
-rw-rw-r-- 1 ec2-user ec2-user 197153999 May 3 18:35 hadoop-2.6.4.tar.gz
-rw-rw-r-- 1 ec2-user ec2-user 146 Oct 1 22:15 lucky.py
-rw-rw-r-- 1 ec2-user ec2-user 298 Oct 1 22:17 lucky.txt
-rw-rw-r-- 1 ec2-user ec2-user 69 Oct 1 19:15 mycopy.txt
-rw-rw-r-- 1 ec2-user ec2-user 102 Oct 1 21:50 myFile.txt
-rw-rw-r-- 1 ec2-user ec2-user 417 Oct 1 20:28 myFile.zip
-rw-rw-r-- 1 ec2-user ec2-user 33 Oct 1 21:05 redirect1.txt
-rw-rw-r-- 1 ec2-user ec2-user 368 Oct 1 21:06 redirect1.txt
-rw-rw-r-- 1 ec2-user ec2-user 48 Oct 1 22:20 was.py
-rw-rw-r-- 1 ec2-user ec2-user 1884 Oct 2 01:00 wordCount.py
[ec2-user@ip-172-31-21-33 ~]$ hadoop fs -put bioproject.xml /data
[ec2-user@ip-172-31-21-33 ~]$ hadoop fs -ls /data
Found 1 items
-rw-rw-r-- 3 ec2-user supergroup 231149003 2021-10-02 03:42 /data/bioproject.xml
```

Run word count on the downloaded text file, using the time command to determine the total runtime of the MapReduce job. You can use the following (single-line!) command. This invokes the wordcount example built into the example jar file, supplying /data/bioproject.xml as the input and /data/wordcount1 as the output directory. Please remember this is one command, if you do not paste it as a single line, it will not work.

```
time hadoop jar hadoop-2.6.4/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.4.jar wordcount /data/bioproject.xml /data/wordcount1
```

Report the time that the job took to execute as screenshot

```
ec2-user@ip-172-31-21-23:~$ cat /data/wordcount1/part-r-00000
21/10/02 03:54:34 INFO mapreduce.Job: map 100% reduce 100%
21/10/02 03:54:34 INFO mapreduce.Job: Job Job_1633145348467_0001 completed successfully
21/10/02 03:54:35 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=54605201
  FILE: Number of bytes written=4627943
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=23113399
  HDFS: Number of bytes written=20056175
  HDFS: Number of read operations=9
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=0
  Launched reduce tasks=1
  Data-local map tasks=0
  Total time spent by all maps in occupied slots (ms)=94173
  Total time spent by all reduces in occupied slots (ms)=12295
  Total time spent by all map tasks (ms)=94173
  Total time spent by all reduce tasks (ms)=12295
  Total vcore-milliseconds taken by all map tasks=94173
  Total vcore-milliseconds taken by all reduce tasks=12295
  Total megabyte-milliseconds taken by all map tasks=46433152
  Total megabyte-milliseconds taken by all reduce tasks=12590080
Map-Reduce Framework
  Map input records=5284546
  Map output records=18562366
  Map output bytes=27956609
  Map output materialized bytes=26902454
  Input split bytes=202
  Combine input records=20053191
  Combine output records=2673165
  Reduce input groups=1040390
  Reduce shuffle bytes=24902454
  Reduce input records=182340
  Reduce output records=1040390
  Spilled Records=1855105
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=1052
  CPU time spent (ms)=1180
  Physical memory (bytes) snapshot=566173696
  Virtual memory (bytes) snapshot=6320132096
  Total committed heap usage (bytes)=334364672
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=23113399
File Output Format Counters
  Bytes Written=20056175
real    1m12.495s
user    0m3.885s
sys     0m0.268s
[ec2-user@ip-172-31-21-23 ~]$
```

(this reports the size of a particular file or directory in HDFS. The output file will be named part-r-00000)

hadoop fs -du /data/wordcount1/

(Just like in Linux, the cat HDFS command will dump the output of the entire file and grep command will filter the output to all lines that matches this particular word). To determine the count of occurrences of “arctic”, run the following command:

hadoop fs -cat /data/wordcount1/part-r-00000 | grep arctic

It outputs the entire content of part-r-00000 file and then uses pipe | operator to filter it through grep (filter) command. If you remove the pipe and grep, you will get the entire word count content dumped to screen, similar to cat command.

Congratulations, you just finished running wordcount using Hadoop.

Part 4: Hive Intro

- 4) In this section we are going to use Hive to run a few queries over the Hadoop framework. These instructions assume that you are starting from a working Hadoop installation. If you are starting your instance, you need to start Hadoop as well. Hive commands are listed in **Calibri bold font**

- a) Download and install Hive:

cd

(this command is there to make sure you start from home directory, on the same level as where hadoop is located)

```
wget http://dbgroup.cdm.depaul.edu/Courses/CSC555/apache-hive-2.0.1-bin.tar.gz  
gunzip apache-hive-2.0.1-bin.tar.gz  
tar xvf apache-hive-2.0.1-bin.tar
```

set the environment variables (can be automated by adding these lines in ~/.bashrc). If you don't, you will have to set these variables every time you use Hive.

```
export HIVE_HOME=/home/ec2-user/apache-hive-2.0.1-bin  
export PATH=$HIVE_HOME/bin:$PATH
```

```
$HADOOP_HOME/bin/hadoop fs -mkdir /tmp  
$HADOOP_HOME/bin/hadoop fs -mkdir /user/hive/warehouse  
(if you get an error here, it means that /user/hive does not exist yet. Fix that by running  
$HADOOP_HOME/bin/hadoop fs -mkdir -p /user/hive/warehouse instead)
```

```
$HADOOP_HOME/bin/hadoop fs -chmod g+w /tmp  
$HADOOP_HOME/bin/hadoop fs -chmod g+w /user/hive/warehouse
```

We are going to use Vehicle data (originally from <http://www.fueleconomy.gov/feg/download.shtml>)

You can get the already unzipped, comma-separated file from here:

```
wget http://dbgroup.cdm.depaul.edu/Courses/CSC555/vehicles.csv
```

You can take a look at the data file by either

```
nano vehicles.csv or
```

```
more vehicles.csv (you can press space to scroll and q or Ctrl-C to break out)
```

Note that the first row in the data is the list of column names. What follows after commands that start Hive, is the table that you will create in Hive loading the first 5 columns. Hive is not particularly sensitive about invalid or partial data, hence if we only define the first 5 columns, it will simply load the first 5 columns and ignore the rest.

You can see the description of all the columns here (atvtype was added later)

<http://www.fueleconomy.gov/feg/ws/index.shtml#vehicle>

Create the ec2-user directory on the HDFS side (absolute path commands should work anywhere and not just in Hadoop directory as bin/hadoop does). Here, we are creating the user "home" directory on the HDFS side.

```
hadoop fs -mkdir /user/ec2-user/
```

Run hive (from the hive directory because of the first command below):

```
cd $HIVE_HOME
```

```
$HIVE_HOME/bin/schematool -initSchema -dbType derby
```

(NOTE: This command initializes the database metastore. If you need to restart/reformat or see errors related to meta store, delete the metastore using **rm -rf metastore_db/** and then repeat the above initSchema command)

bin/hive

You can now create a table by pasting this into the Hive terminal:

```
CREATE TABLE VehicleData (  
barrels08 FLOAT, barrelsA08 FLOAT,  
charge120 FLOAT, charge240 FLOAT,  
city08 FLOAT)  
ROW FORMAT DELIMITED FIELDS  
TERMINATED BY ',' STORED AS TEXTFILE;
```

You can load the data (from the local file system, not HDFS) using:

```
LOAD DATA LOCAL INPATH '/home/ec2-user/vehicles.csv'  
OVERWRITE INTO TABLE VehicleData;
```

(NOTE: If you downloaded vehicles.csv file into the hive directory, you have to change file name to /home/ec2-user/apache-hive-2.0.1-bin/vehicles.csv instead)

Verify that your table had successfully loaded by running

```
SELECT COUNT(*) FROM VehicleData;
```

(Copy the query output and report how many rows you got as an answer.)

```
hive> SELECT COUNT(*) FROM VehicleData;  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.  
Query ID = ec2-user_20211002233719_be99c358-75d2-4319-9868-57e8fee15da4  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks determined at compile time: 1  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reducers=<number>  
Starting Job = job_1633211013371_0001, Tracking URL = http://ip-172-31-21-33.us-east-2.compute.internal:8088/proxy/application_1633211013371_0001/  
Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job_1633211013371_0001  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1  
2021-10-02 23:37:33,303 Stage-1 map = 0%, reduce = 0%  
2021-10-02 23:37:41,139 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.32 sec  
2021-10-02 23:37:50,027 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.62 sec  
MapReduce Total cumulative CPU time: 2 seconds 620 msec  
Ended Job = job_1633211013371_0001  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.62 sec HDFS Read: 11775003 HDFS Write: 6 SUCCESS  
Total MapReduce CPU Time Spent: 2 seconds 620 msec  
OK  
34175  
Time taken: 32.548 seconds, Fetched: 1 row(s)  
hive>
```

There are 34,175 rows.

Run a couple of HiveQL queries to verify that everything is working properly:

SELECT MIN(barrels08), AVG(barrels08), MAX(barrels08) FROM VehicleData;
(copy the output from that query)

```
hive> SELECT MIN(barrels08), AVG(barrels08), MAX(barrels08) FROM VehicleData;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = ec2-user_20211002234324_889f978a-d5cc-4b49-a540-e8210692aa35
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1633211013371_0002, Tracking URL = http://ip-172-31-21-33.us-east-2.compute.internal:8088/proxy/application_1633211013371_0002/
Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job_1633211013371_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-10-02 23:43:32,234 Stage-1 map = 0%, reduce = 0%
2021-10-02 23:43:40,009 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.94 sec
2021-10-02 23:43:46,610 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.25 sec
MapReduce Total cumulative CPU time: 3 seconds 250 msec
Ended Job = job_1633211013371_0002
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.25 sec HDFS Read: 11777415 HDFS Write: 37 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 250 msec
OK
0.059852      17.820177449476272      47.06831
Time taken: 25.583 seconds, Fetched: 1 row(s)
hive>
```

As shown above:

- The Minimum of barrels08 is 0.06
- The Average of barrels08 is 17.82
- The Maximum of barrels08 is 47.07

SELECT (barrels08/city08) FROM VehicleData;
(you do not need to report the output from that query, but report “Time taken”)

The time taken is 0.194 seconds to fetch 34,175 rows as shown below:

```
0.6539881115867978
0.4898370901743571
0.6278285526093983
0.7875944438733553
0.7163524150848388
0.8716353310479058
0.8716353310479058
1.1440213918685913
Time taken: 0.194 seconds, Fetched: 34175 row(s)
hive>
```

Next, we are going to output three of the columns into a separate file (as a way to transform data for further manipulation that you may be interested in)

INSERT OVERWRITE DIRECTORY 'ThreeColExtract'
SELECT barrels08, city08, charge120
FROM VehicleData;

You can now exit Hive by running **exit**;

And verify that the new output file has been created (the file will be called 000000_0)

The file would be created in HDFS in user home directory (/user/ec2-user/ThreeColExtract)

Report the size of the newly created file and include the screenshot.

The size is 613.2 K as shown below:

```
ec2-user@ip-172-31-21-33:~/apache-hive-2.0.1-bin
[ec2-user@ip-172-31-21-33 apache-hive-2.0.1-bin]$ hadoop fs -ls -h /user/ec2-user/ThreeColExtract
Found 1 items
-rwxr-xr-x   3 ec2-user supergroup    613.2 K 2021-10-02 23:58 /user/ec2-user/ThreeColExtract/000000_0
[ec2-user@ip-172-31-21-33 apache-hive-2.0.1-bin]$
```

Next, you should go back to the Hive terminal, create a new table that is going to load 8 columns instead of 5 in our example (i.e. create and load a new table that defines 8 columns by including columns city08U,cityA08,cityA08U) and use Hive to generate a new output file containing only the city08U and cityA08U columns from the vehicles.csv file. Report the size of that output file as well.

If the size is zero, you are looking at a directory and the output file(s) are in that directory.

The size of the file is 281.0K as shown below:

```
ec2-user@ip-172-31-21-33:~/apache-hive-2.0.1-bin
[ec2-user@ip-172-31-21-33 apache-hive-2.0.1-bin]$ hadoop fs -ls -h /user/ec2-user/TwoColExtract
Found 1 items
-rwxr-xr-x   3 ec2-user supergroup    281.0 K 2021-10-03 00:28 /user/ec2-user/TwoColExtract/000000_0
[ec2-user@ip-172-31-21-33 apache-hive-2.0.1-bin]$
```

Submit a single document containing your written answers. Be sure that this document contains your name and “CSC 555 Assignment 2” at the top.