

# CSC 555 and DSC 333

## Mining Big Data

### Lecture 2

Alexander Rasin

College of CDM, DePaul University

September 21<sup>st</sup>, 2021

# Tonight


- MapReduce/Hadoop
  - Distributed computing
  - Performance considerations
- Hashing
- SQL to MapReduce
- Hive

# Distributed Computing



# Parallel Computing is Hard

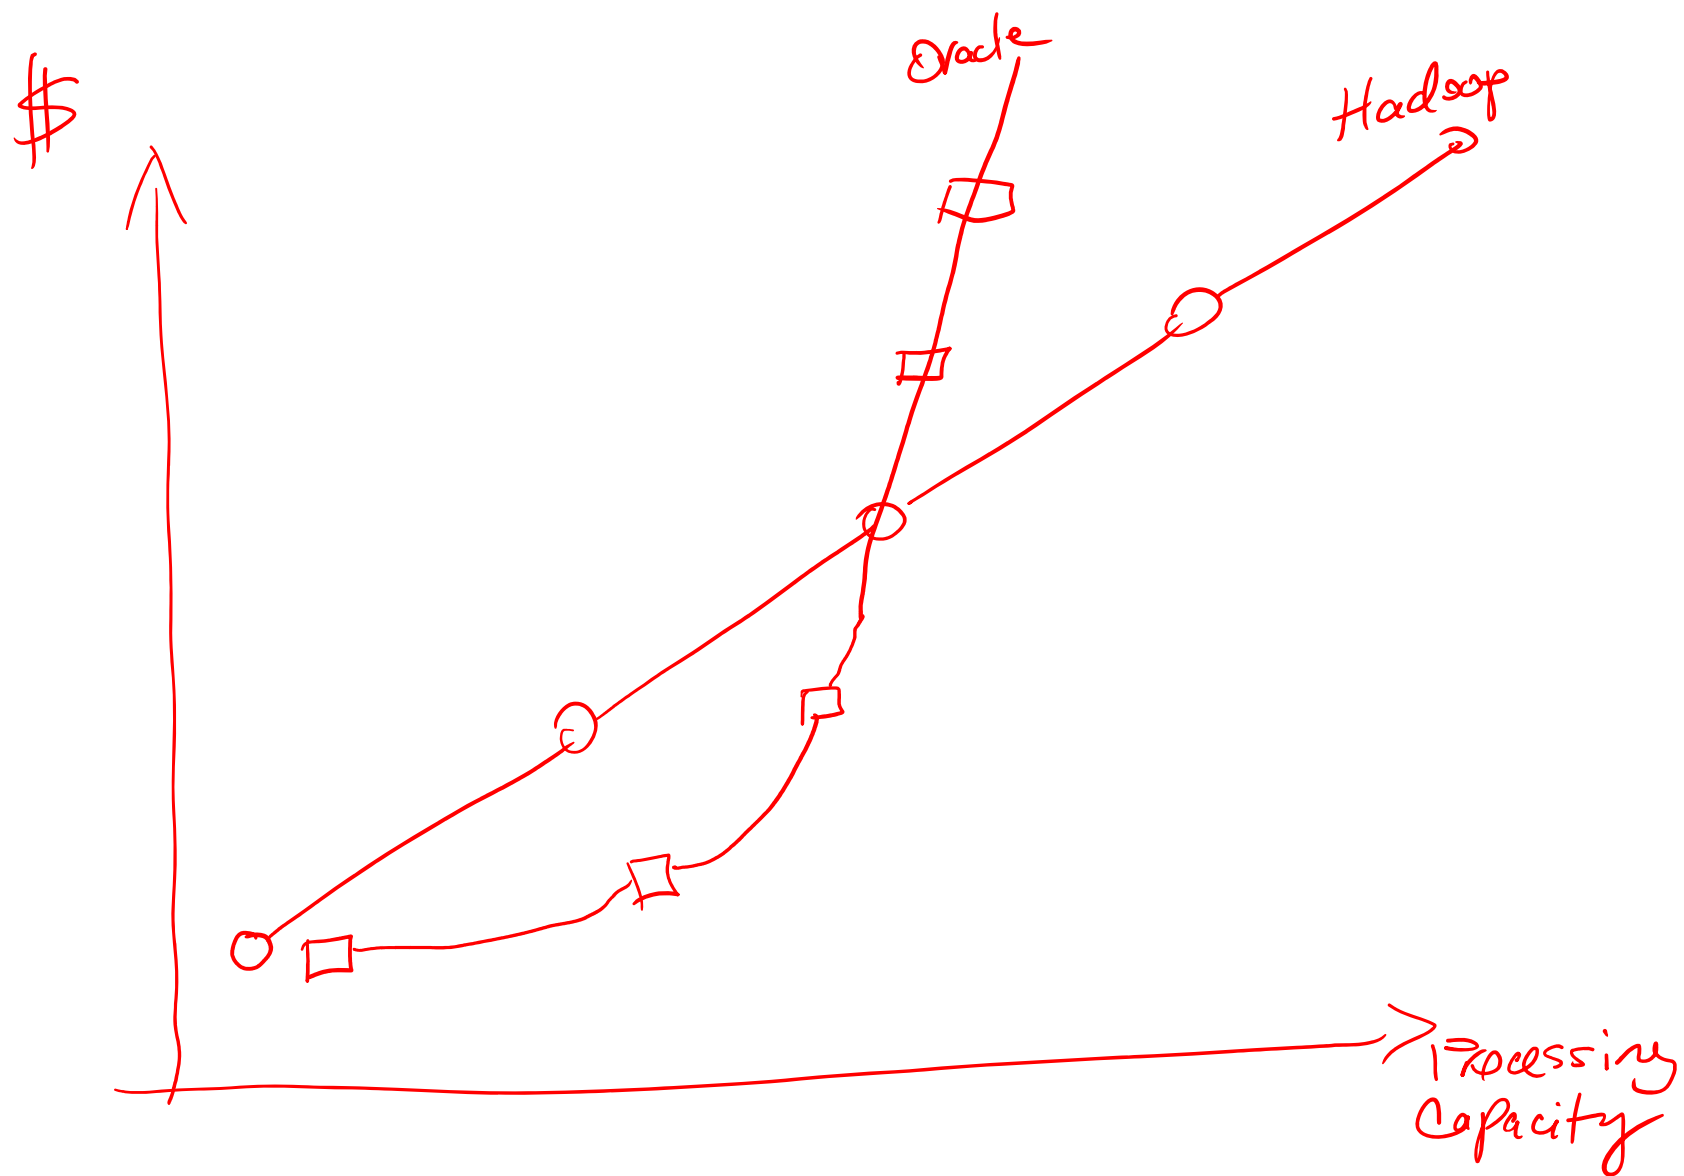
- Fundamental issues



# Scale-out vs Scale-up


- Many smaller servers
- Harder to manage
- Easier to expand
- Easier to tolerate failure (\*)
- Generally cheaper

- Few large servers
- Easier to manage
- Difficult/expensive to expand
- Failure is a serious setback
- More expensive



# Hadoop Execution

- Every job represented as
  - Map function
  - (Optional) Combine function
  - Reduce function
- Provides convenient failure/retry semantics
- Extensive Ecosystem
  - Hive, Pig, Spark, Storm, HBase, Mahout, Zookeeper, Sqoop, HUE, Oozie, Avro, Flume, ...

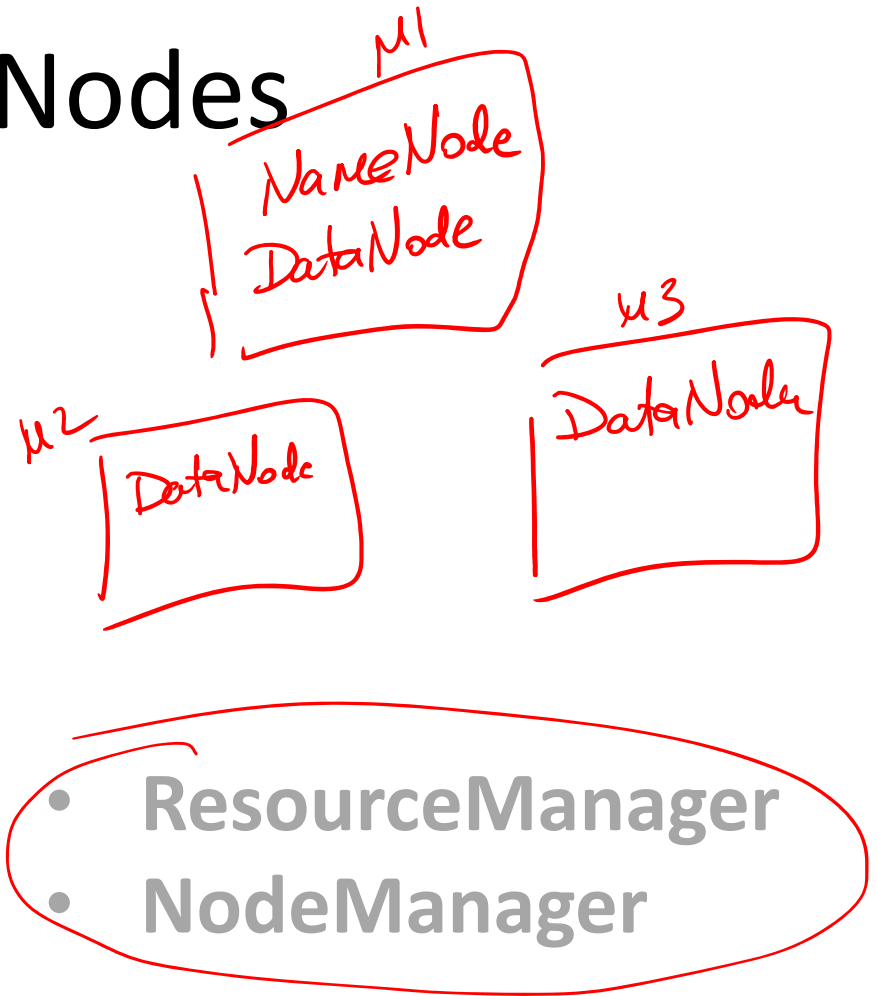


# Types of Nodes

- HDFS Nodes

- NameNode
- DataNode

*Supervisor*



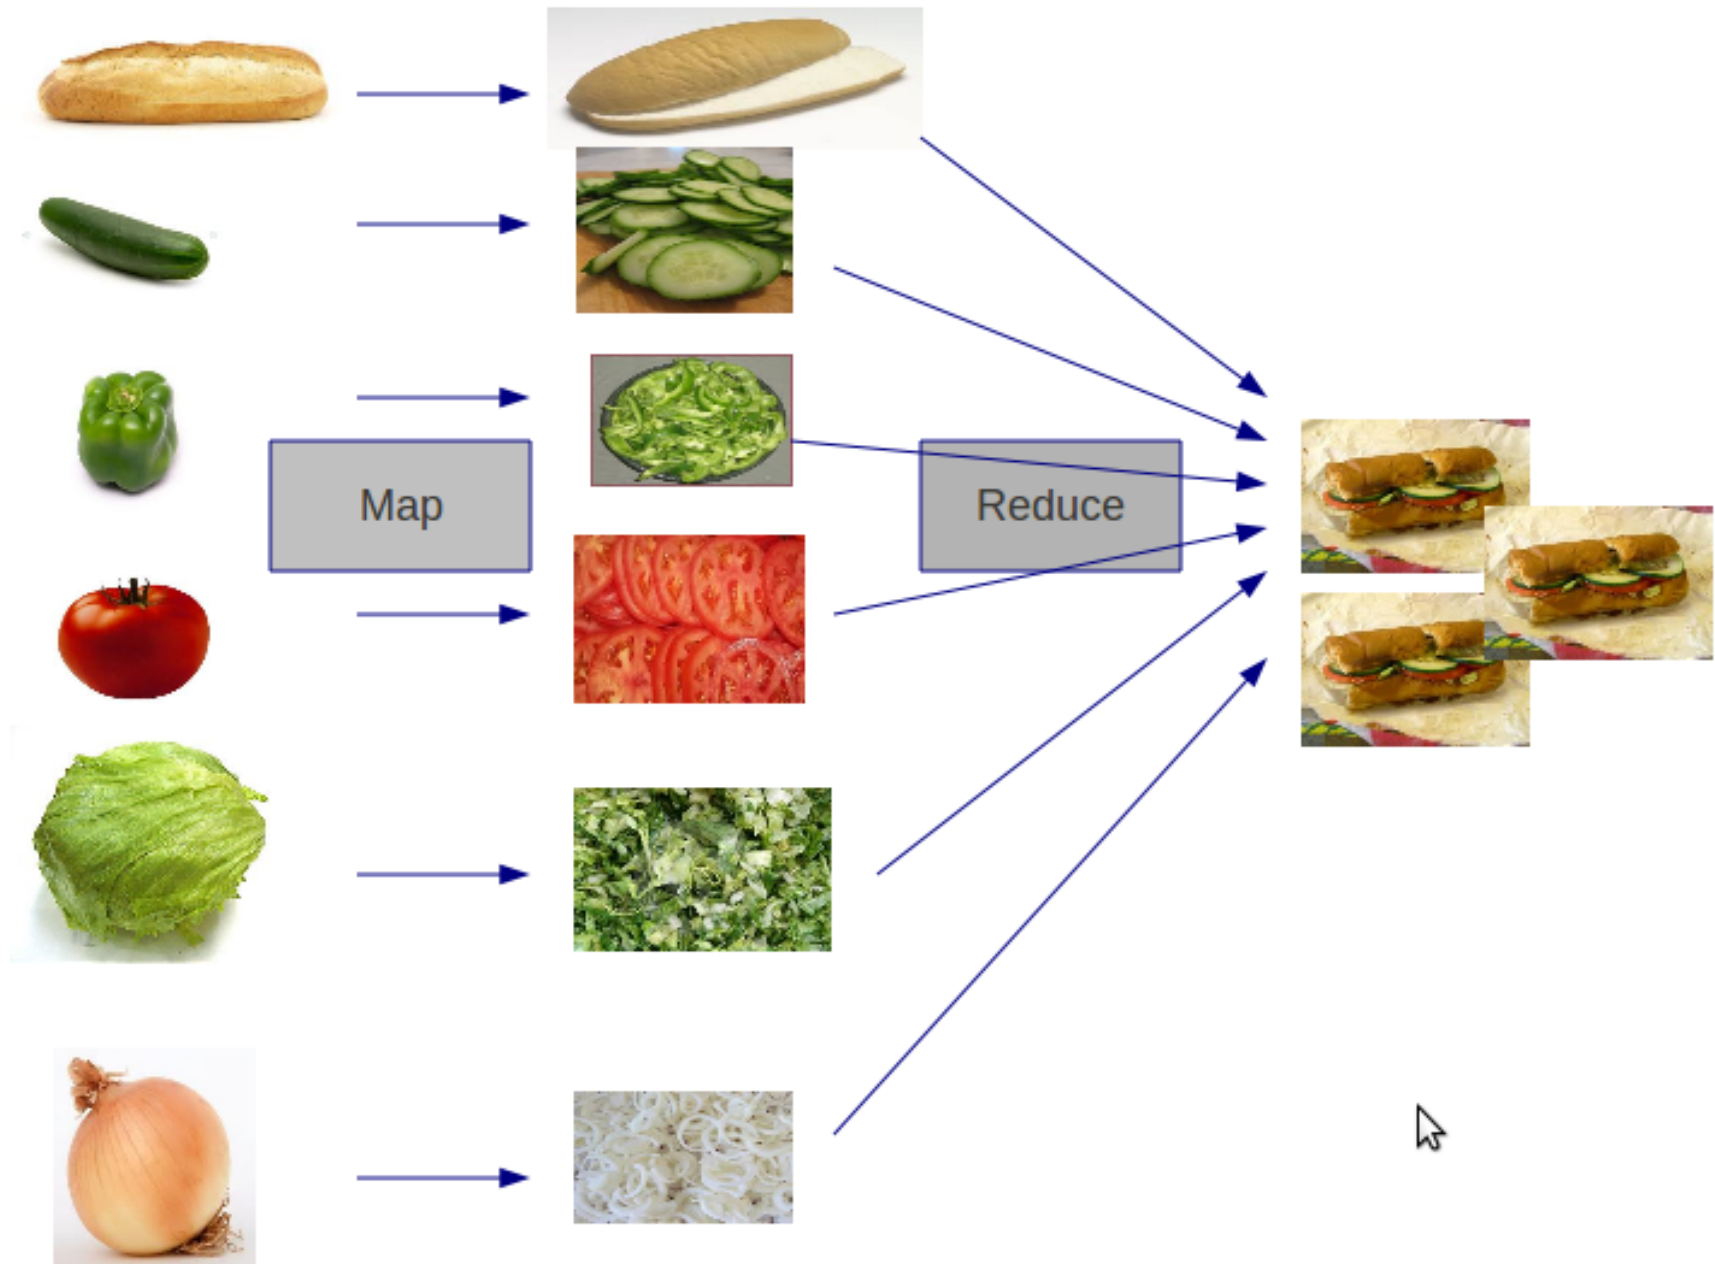
- MapReduce Nodes

- JobTracker
- TaskTracker

- ResourceManager
- NodeManager

- Other nodes (checkpoint, balancer, etc)





# Counting Words

select word, count(\*)  
FROM myWords  
Group BY word

INPUT

this is an example text  
this is not an example  
this text text

set of  
blocks

B1  
□

B2  
✓

B3  
✓

OUTPUT

an, 2  
example, 2  
is, 2  
not, 1  
text, 3  
this, 3

B3


M3


this 1  
text 1  
text 1

- Applied on **per-block** basis
- Assigned to a machine

# Map Worker

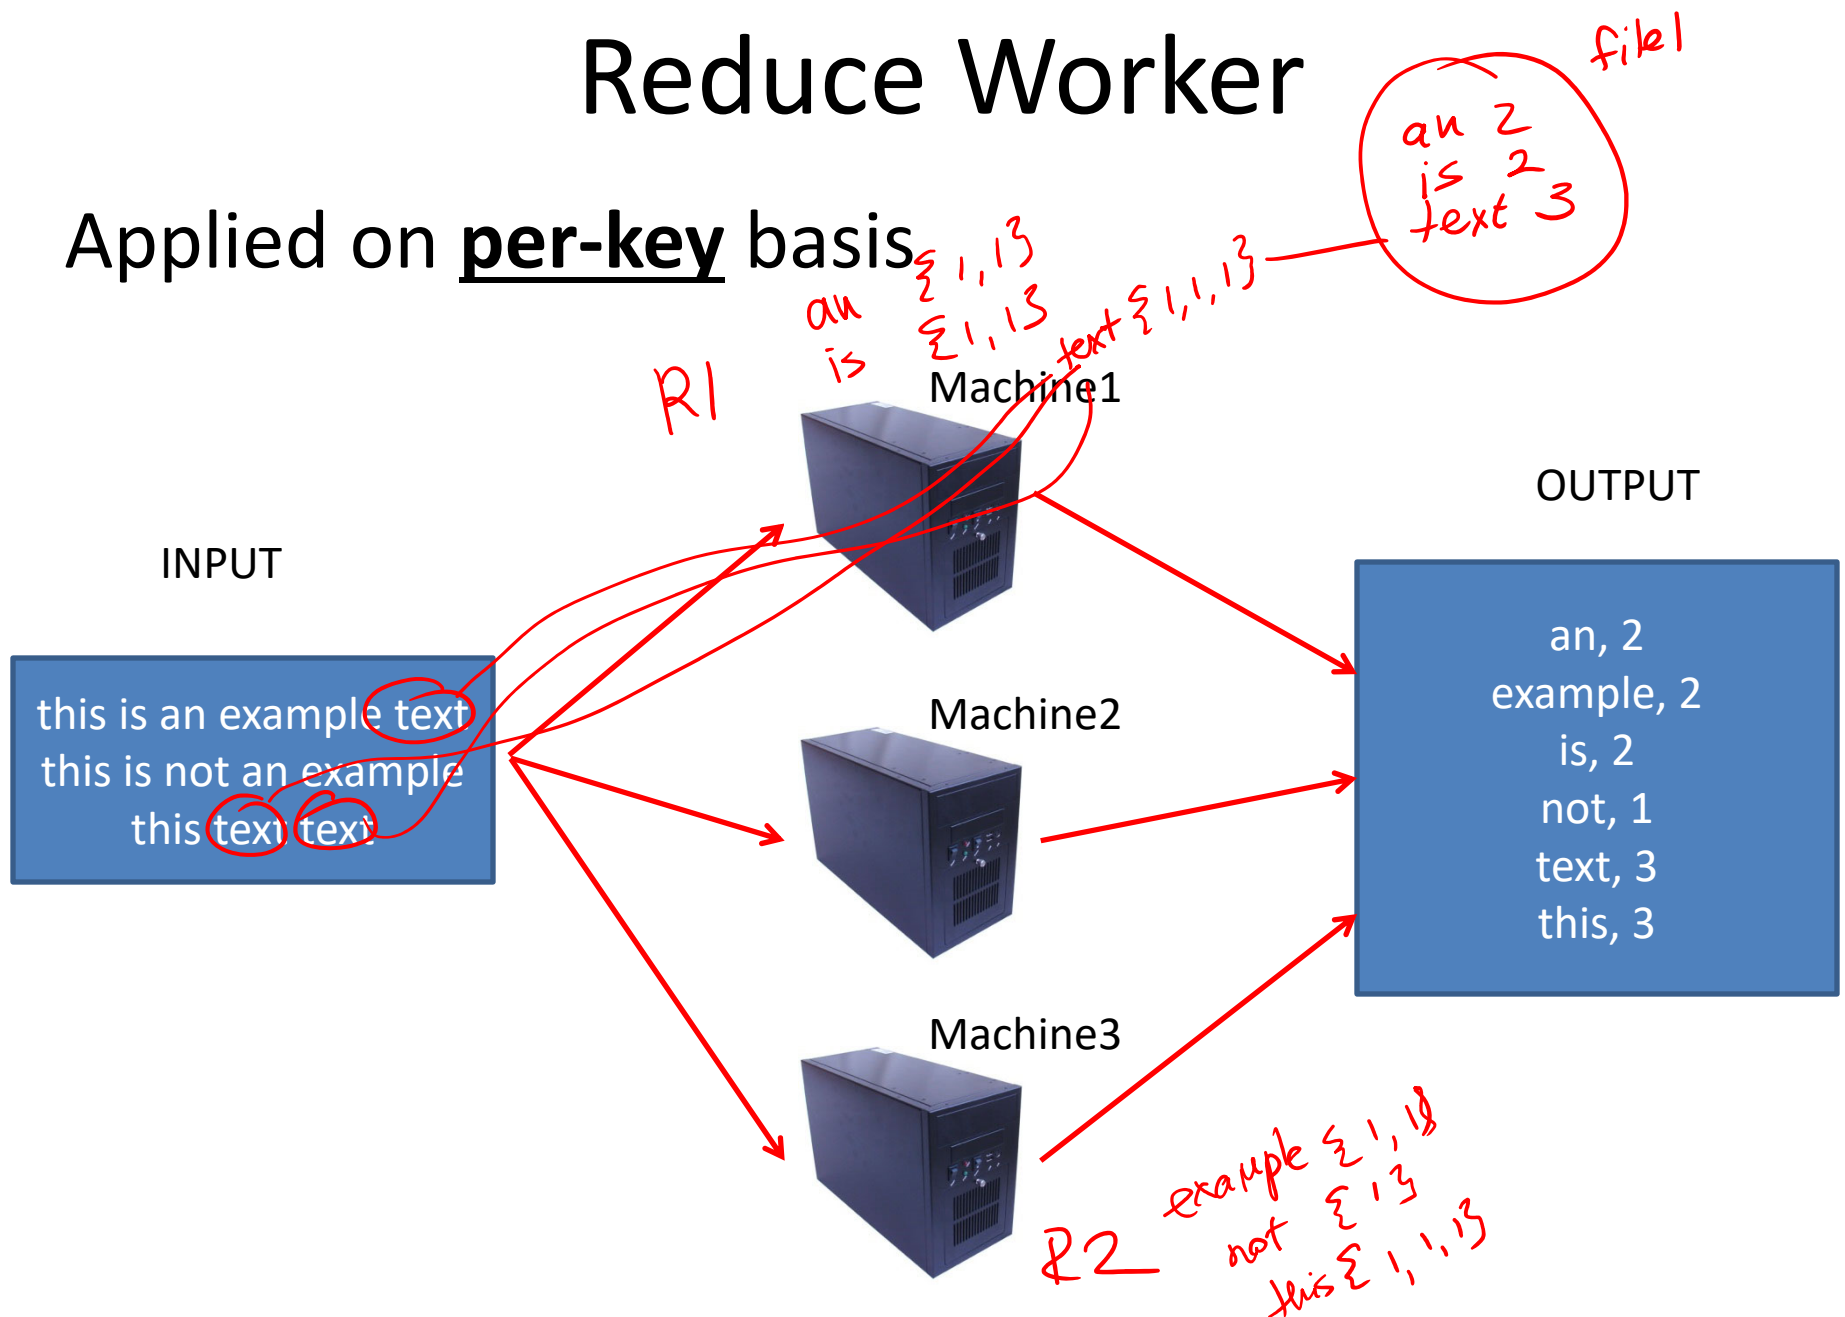
Applied on per-block basis






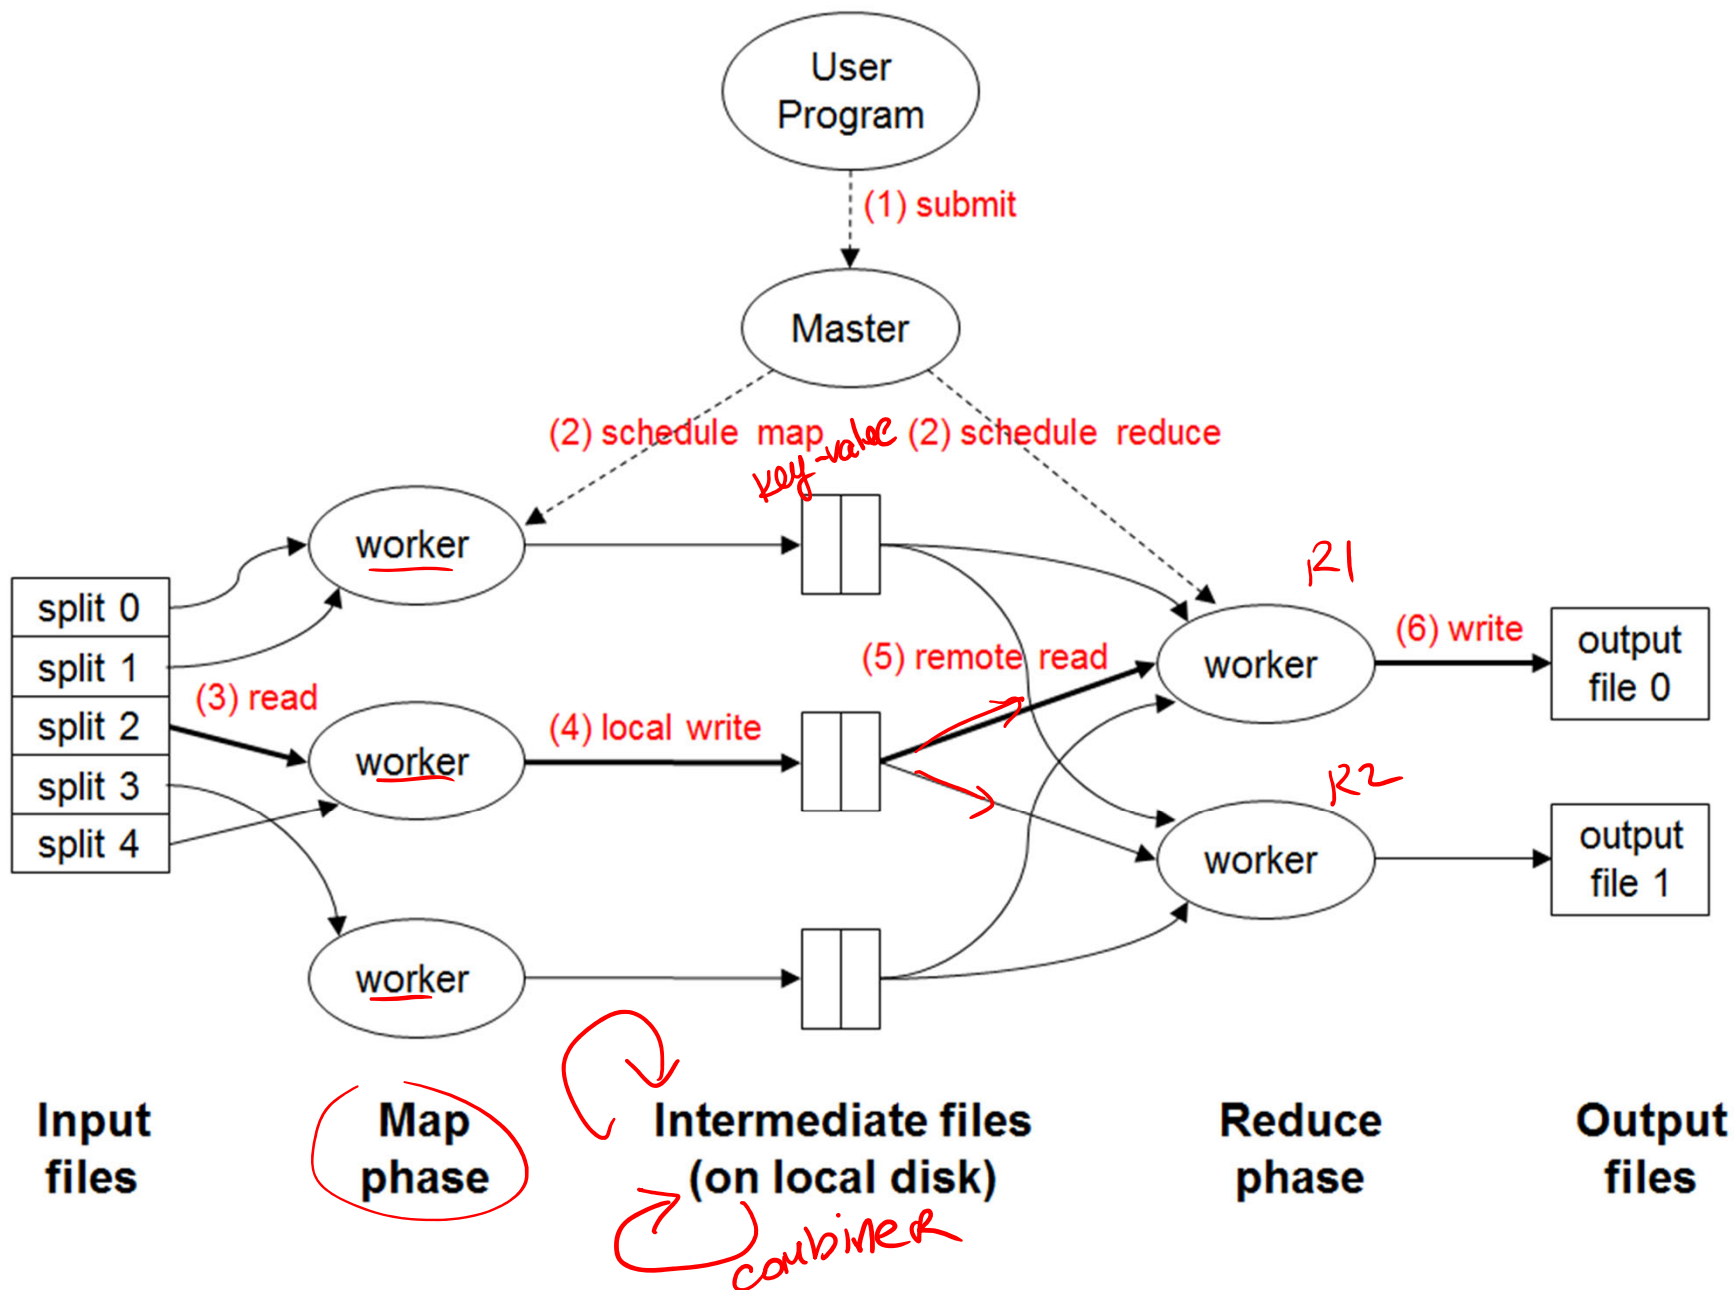
# Reduce Worker

Applied on per-key basis



# Distributed Word Count







# Java Code: Map

```
/**
 * Counts the words in each line.
 * For each line of input, break the line into words and emit them as
 * (<b>word</b>, <b>1</b>).
 */
public static class MapClass extends MapReduceBase
    implements Mapper<LongWritable, Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value,
        OutputCollector<Text, IntWritable> output,
        Reporter reporter) throws IOException {
        String line = value.toString();
        StringTokenizer itr = new StringTokenizer(line);
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            output.collect(word, one);
        }
    }
}
```

*line.split(' ')*

# Python Code: Map

```
# input comes from STDIN (standard input)
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        #
        # tab-delimited; the trivial word count is 1
        print '%s\t%s' % (word, 1)
```

*Handwritten notes:*

- `sys.stdin:` is circled in red.
- `line.strip()` is underlined in red.
- `line.split()` is underlined in red.
- `['this', 'is', 'text']` is written in red next to `line.split()`.
- `'this is text'` is written in red above `line.strip()`.
- `print '%s\t%s' % (word, 1)` is underlined in red.
- `\t` is circled in red.
- `→ this \t 1  
is \t 1  
text \t 1` is written in red next to the print statement.

# Java Code: Reduce

```
/**
 * A reducer class that just emits the sum of the input values.
 */
public static class Reduce extends MapReduceBase
    implements Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterator<IntWritable> values,
        OutputCollector<Text, IntWritable> output,
        Reporter reporter) throws IOException {
        int sum = 0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        output.collect(key, new IntWritable(sum));
    }
}
```

```
# input comes from STDIN
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()

    # parse the input we got from mapper.py
    word, count = line.split('\t', 1)

    # convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # count was not a number, so silently
        # ignore/discard this line
        continue

    # this IF-switch only works because Hadoop sorts map output
    # by key (here: word) before it is passed to the reducer
    if current_word == word:
        current_count += count
    else:
        if current_word:
            # write result to STDOUT
            print '%s\t%s' % (current_word, current_count)
        current_count = count
        current_word = word
```

# Performance in Hadoop

- Run WordCount with different settings
- Use 1-node and 5-node Hadoop cluster
- File inputs
  - Small (0.5MB), Medium (182MB), Large (5.6GB)
- Results
  - Small (1-node : 0.5min    5-nodes : 0.5min)
  - Medium (1-node : 1.5min    5-nodes : 1min )
  - Large (1-node : 26 min    5-nodes : <8min )

128 154



# A Break

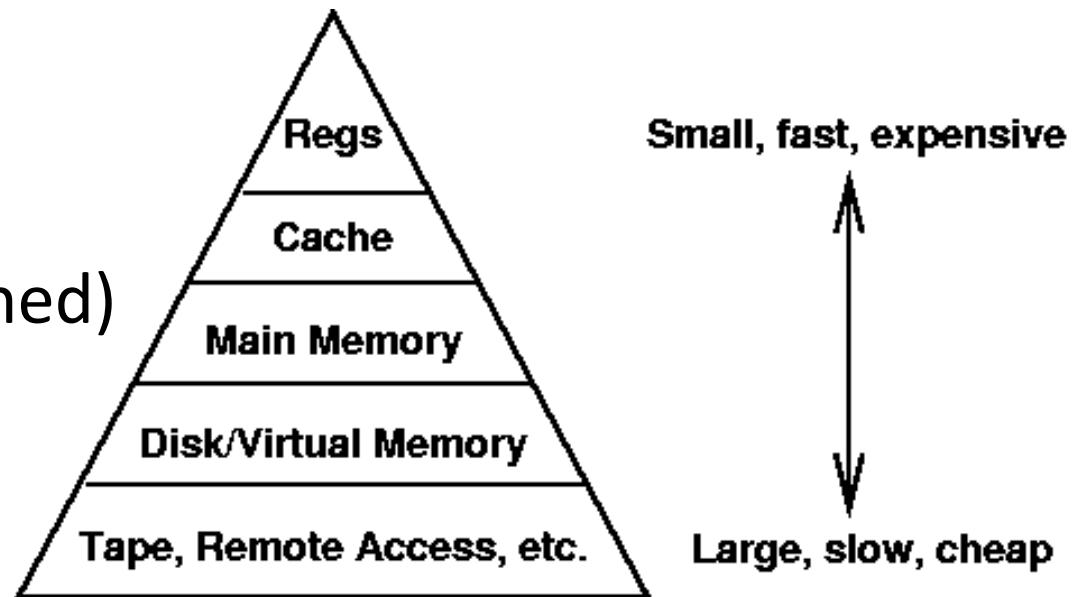


# Runtime / Cost

- Accessing the disk
  - Reading from disk is expensive
- Communication cost
  - Network slower than disk
  - Rack to Rack (Data center to data center)
- Execution cost
  - Actual tasks (map, reduce, etc.)

# Disk Access

- Disks are cheap but slow
  - Spark (in RAM)
  - Compression
- Seek vs Read
  - Small files (batched)
  - In place updates





# Network Transfer

- Transferring data may be slow
  - Local access/disk is faster
  - Computation is much faster
  - Distance (neighbor, other rack, other datacenter)
- Network saturation
  - Same pipe / joins
  - $N^2$  communication cost

# File Balance in HDFS

- Where do map tasks read the files?
- Balance the workload
  - Worst case => entire file is on one node
- Why unbalanced?
  - New (or replacement) nodes are added
  - Extra drives added, heterogeneous machines

# How is HDFS (Map) balanced?

- Blocks written to disk at random
  - (hopefully)
- No automatic rebalancing
  - Replace file (i.e. write again)
  - Turn replication off/on (rewrite duplicates)
  - Run a rebalancer

# Reduce Balancing

- How many values does each reducer process?
- Values are assigned across reducers
  - Hashing function (default or custom)
- Worst case => all keys go to just one reducer

# What is Hashing?

- The **Hash House Harriers** (abbreviated to **HHH** or **H3**, or referred to simply as **hashing**) is an international group of non-competitive running social clubs. An event organized by a club is known as a **hash** or **hash run**, with participants calling themselves **hashers** or **hares, hounds, harriers, and harriets**.

# HISTORY OF THE HASH




May 23, 1874 ]

HARE AND HOUNDS  
HARPER'S WEEKLY

44 ]

# Hashing

- **$h(k) \text{ MOD } N$**  ( $N = \#$  of buckets)
  - bucket to which data entry with key  $k$  belongs.



# Hashing

- Buckets correspond to Reducers
- Hash function works on **key** field of record **r**.  
Use its value MOD N to distribute values over range 0 ... N-1.
  - $h(\text{key}) = (a * \text{key} + b)$  usually works well
  - a and b are constants
- Node may have different number of reducers



# Hadoop Partitioner

```
public class HashPartitioner<K, V> extends  
Partitioner<K, V> {  
    /** Use {@link Object#hashCode()} to partition. */
```

```
    public int getPartition(K key, V value,  
                           int numReduceTasks) {  
        return (key.hashCode() & Integer.MAX_VALUE) %  
numReduceTasks;  
    }  
}
```

*MOD*

# Sorting Hat



# Describing MapReduce algorithms

- Initially pseudo-code like
  - Later we will use python
- Discuss examples for Map and Reduce tasks
  - Describe the tasks
  - What does Map do?
    - Per block: identify the key, identify the value
  - What does Reduce do?
    - Per key: apply the function

# Queries with MapReduce

SELECT \* FROM Students

- Hadoop
  - Parses the input file every time
  - Applies filter when reading/parsing

SELECT \* FROM Students  
WHERE age > 30 AND Year < 2022;

Map      <sup>key</sup> ID      <sup>value</sup> ~~other~~ Columns

print( ID + '/' + restofrow )  
if year < 2022 :  
    print( ID + '/' + restofrow )

# Queries with MapReduce

SELECT Age, COUNT(\*)

FROM Student

GROUP BY Age

K1 31 {A, B, C, A, B} R2 32 {A, A, A, A, A}  
33 {C, C, C, C, C}

Avg(Grade)

Map

Age

Grade

Reducer

Age

Avg(Grade)

SELECT Age, Year COUNT(\*)

FROM Student

GROUP BY Age, Year

31 - 2022 1  
31 - 2023

Map

key  
Age-Year

Value

1

Reduce

Age, Year

SUM(1)

# Set Union

- MapReduce has to parse both files

- SQL:

```
SELECT ID FROM Student
UNION
SELECT ID FROM Faculty;
```

	key	Value
<sup>S</sup> Map1	ID <sub>S</sub>	1
<sup>F</sup> Map2	ID <sub>F</sub>	1
Reducer	ID	—

# Union with MapReduce

- Need to process 2 files
  - Input == directory with 2 files
- Map
  - Process each file
  - Emit compatible keys
- Reduce
  - Emit single output for every key

# Set Intersection

$(7) \cap (5)$   
 $7S$   
 $7F$

- MapReduce has to parse both files

	key	Value
Map 1	IDs	<del>7</del> S
Map 2	IDs	<del>7</del> F

- SQL:

```
SELECT ID FROM Students
INTERSECT
SELECT ID FROM Faculty;
```

Reducer

7	<del>{1, 3}</del>	{S}
8	{1, 13}	{S, S, F}
9	{1, 13}	{F, F}



# Intersection with MapReduce

- Process two input files
- Need to identify which key is from where
  - Two different Mappers
- Reduce
  - Iterate through keys
  - Check if at least one from each file appears
  - Emit that key
- Does file order matter?

# Hive Architecture



Student.txt

ID	Name	Year
1	Alex	2015
2	Jane	2016
3	Jane	2017

# Hive Execution Flow

- Parse the query
- Get metadata from MetaStore
- Create a logical plan
- Optimize the plan
- Create a physical plan
  - DAG of MR jobs

# Hive Examples

```
CREATE TABLE u_data ( userid INT,  
movieid INT,  
rating INT,  
unixtime STRING)  
ROW FORMAT DELIMITED FIELDS  
TERMINATED BY '\t' STORED AS TEXTFILE; (not compressed)
```

- show tables; describe u\_data;
- wget <http://www.grouplens.org/system/files/ml-100k.zip>

```
LOAD DATA LOCAL INPATH 'ml-100k/u.data'  
OVERWRITE INTO TABLE u_data;
```

# Using Hive

- `SELECT COUNT(*) FROM u_data;`
- `SELECT * FROM u_data WHERE userid = 449;`
- `SELECT userid, AVG(rating) from u_data  
GROUP BY userid;`
- `SELECT userid, AVG(rating) as AR from u_data  
GROUP BY userid ORDER BY AR;`

# Next Time:

- Distributed system/Hadoop performance
- Compression and public-private keys
- Hive
- Read:
  - Mining of Massive Datasets
    - Section 2.3
  - Hadoop: The Definitive Guide
    - pp 3-10, Chapter 1: Data! Through Grid Computing
    - Pp 43-49, Chapter 3: Design of HDFS through HDFS Federation
    - Pp 50-53 Chapter 3: Command-line Interface and Basic Filesystem Operations
    - Pp 69- 71 Chapter 3: Data flow