

# CSC 555 and DSC 333

## Mining Big Data

### Lecture 3

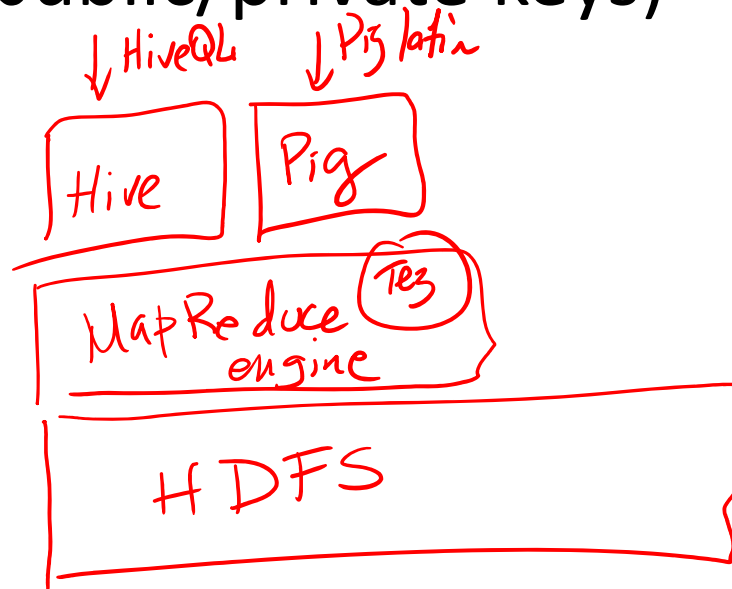
Alexander Rasin

College of CDM, DePaul University

September 28<sup>th</sup>, 2021

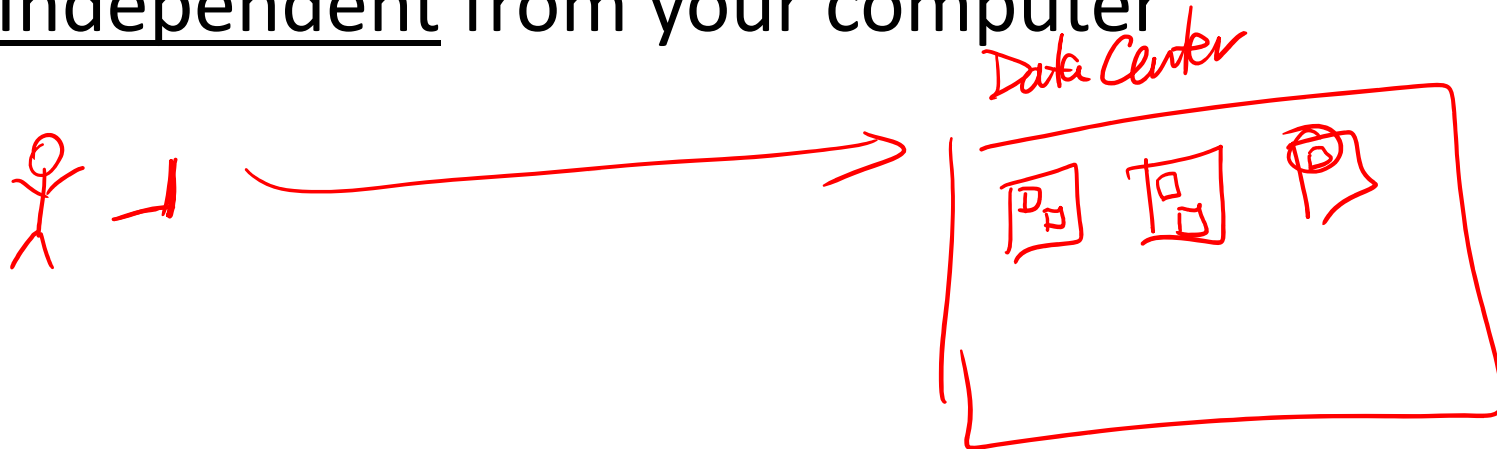
# Tonight

- Virtual instances
- MapReduce
- Cryptography (public/private keys)
- Hive
- Pig



# Cloud Service

- A virtual machine
  - Simulated environment
  - Guaranteed performance
  - On a shared (larger) machine
- Independent from your computer

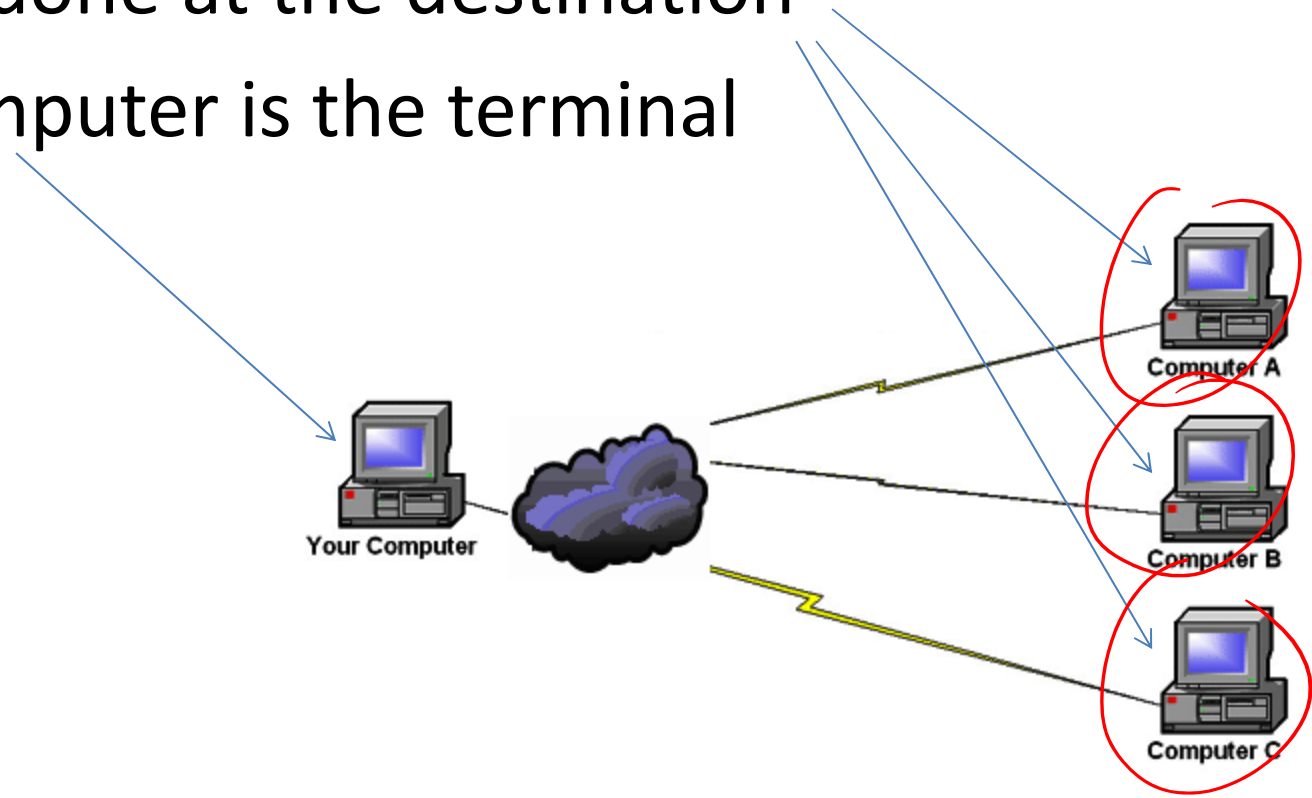


# Hadoop/MapReduce

- Many loosely connected computers
- Limited centralized control
- Therefore
  - Mapper only sees a block at-a-time
  - Reducer only sees its keys
  - Reducer emits one “result” per key

# Remote Access

- Remote connection (such as remote desktop)
- Work is done at the destination
- Your computer is the terminal



# Host Names/IP

- Nodes addressed by
  - Internet Protocol (IP) Address
    - 140.192.5.61
  - Domain Name
    - www.depaul.edu
- IP and Domain name are interchangeable

# Conventions

- localhost = “This machine”
- Hosts
  - ec2-54-187-35-9.us-west-2.compute.amazonaws.com
  - 54.187.35.9

~~54-187-35-9~~

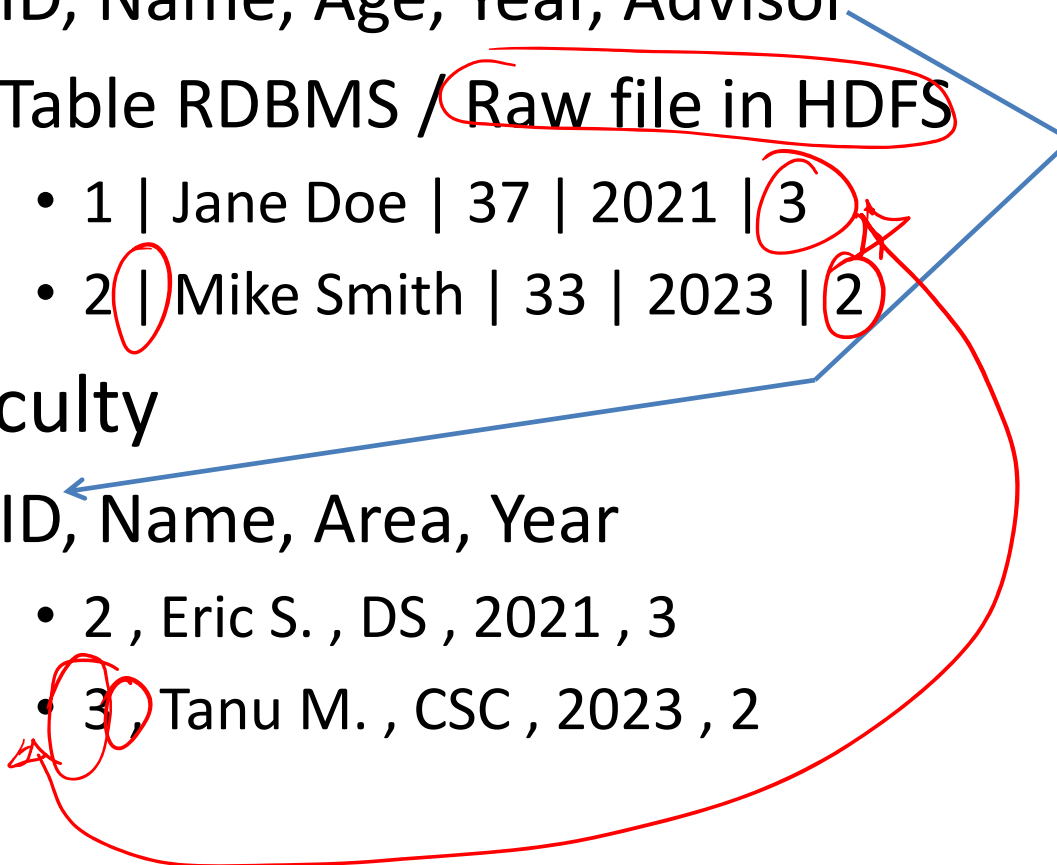
# Student / Faculty

- Student

- ID, Name, Age, Year, Advisor
- Table RDBMS / Raw file in HDFS
  - 1 | Jane Doe | 37 | 2021 | 3
  - 2 | Mike Smith | 33 | 2023 | 2

- Faculty

- ID, Name, Area, Year
  - 2 , Eric S. , DS , 2021 , 3
  - 3 , Tanu M. , CSC , 2023 , 2





# Set Difference (MINUS)

- MapReduce has to parse both files

if line.count('1') == 4 :  
execute Mapper1S code  
else:  
execute Mapper2F code

- SQL:

```
SELECT ID FROM Student  
MINUS (or EXCEPT)  
SELECT ID FROM Faculty;
```

Mapper1S ID S\_L  
Mapper2F ID F\_R  
Reducer ID

X 5 { S\_L, S\_L, F\_R }  
✓ 6 { S\_L, S\_L, S\_L }  
X 7 { F\_R, F\_R }

# Difference

- Process two input files
- Need to identify which key is from where
- Reduce
  - Iterate through all keys
  - Check that only left side key appears
    - Emit output
- Does order matter?

# SQL Join

Mapper 1 S  
Mapper 2 F  
Reducer

Key  
Advisor  
ID  
ID/Advisor  
Name - S  
Area - F  
Name + Area

- SQL:

```
SELECT Student.Name, Faculty.Area
FROM Student, Faculty
WHERE Student.Advisor = Faculty.ID;
```

Alex-S, CSC-F

5 { Alex, CSC }

6 { DSC, Math, CSC }

DSC-F, Math-S, CSC-F

DSC-F, '-S', 'NULL-F'

- Or:

```
SELECT Student.Name, Faculty.Area
FROM Student JOIN Faculty ON Advisor = Faculty.ID;
```

# Join by MapReduce

- Mapper
  - Process both files
  - Identify source of each row
- Reducer (for same key)
  - Collect all values from left file
  - Collect all values from right file
  - Return all (matching) combinations

# Generalizing Join

- Join by 2 attributes (A, B)
  - Rewrite both Mappers to use (A+B) as key
  - Leave Mappers as is and compare/join by the 2<sup>nd</sup> attribute in the Reduce function
- Join without equality? (non-equi-join)
  - E.g., (Student.Age > Faculty.Age)

*S.First = F.First AND*

*S.Last = F.Last*

*First\_Last*

# Grouping/Aggregation

- Key-based aggregation
  - MR join is also similar to an aggregate
- Mapper
  - Same as for join
- Reduce
  - Collect all values
  - Perform aggregation or join

# SQL Join + Aggregation

- Join + Aggregate:

```
SELECT Faculty.Name, COUNT(*)  
FROM Student, Faculty  
WHERE Student.Advisor = Faculty.ID;  
GROUP BY Faculty.Name;
```

Mapper 1s  
Mapper 2F  
Reducer 1

Advisor - S
ID Name
Tanu
Tanu
Tanu
Jay
Jay

↓  
Tanu  
Tanu  
Tanu  
Jay  
Jay

Mapper 3  
Reducer

key	Value
Name	-
Name	Count(*)

↓  
Tanu 3  
Jay 2

# Complex Queries

- N-way joins
- ~~Non-equi join~~
- Complex predicates
- Complex aggregation



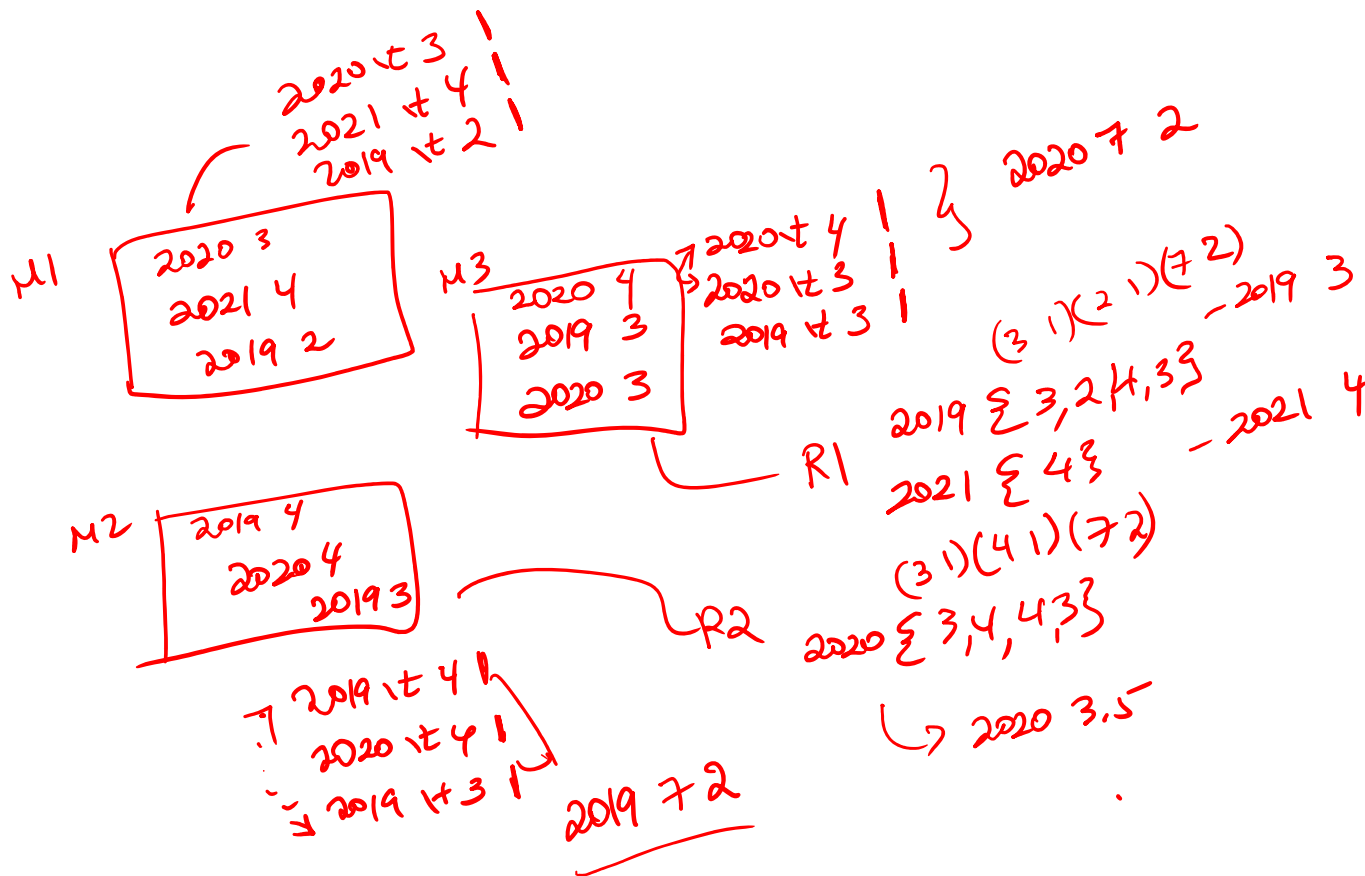
# Computing GPA (by year)

*Select Year, AVG(Grade)  
FROM Records  
GROUP BY Year*

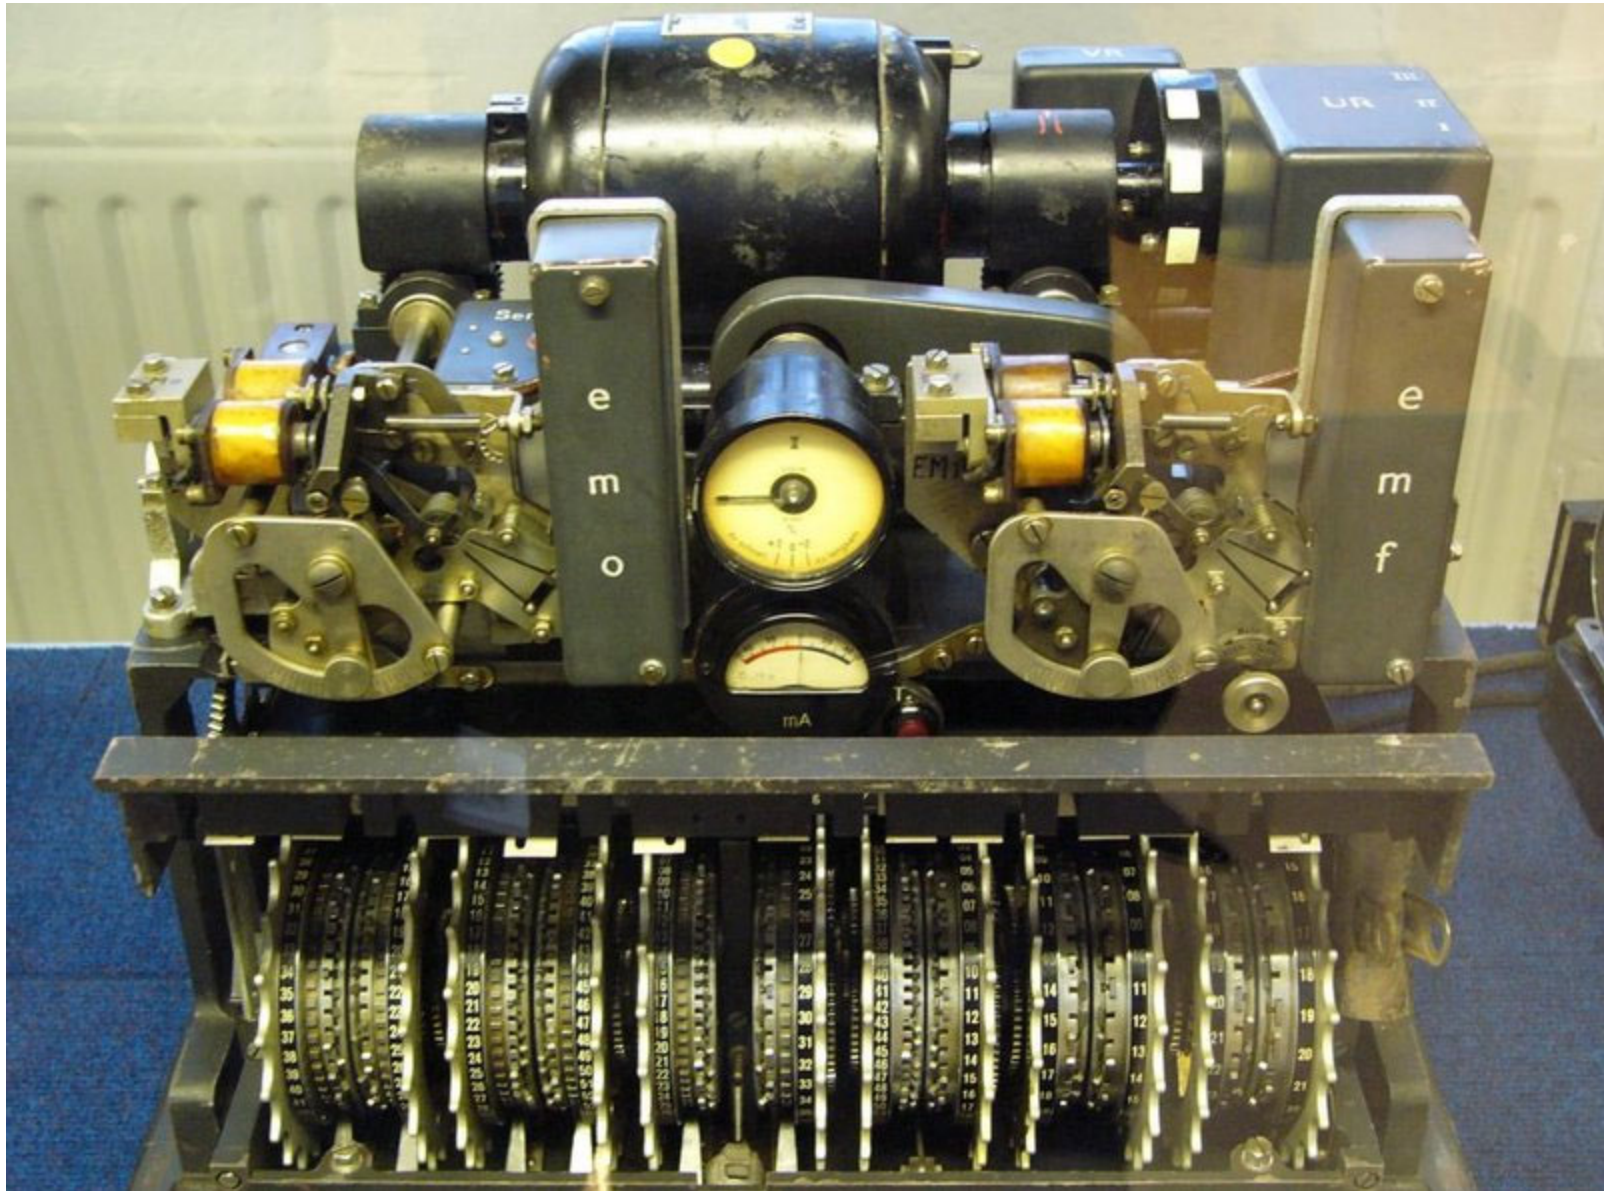
- Map
  - Read/parse the file
  - Emit pairs (year, grade)
- Reduce
  - Sum up all the enrollment values
  - Divide by the count of values
  - Produce (year, avg\_grade) output
- Combiner?

	key	value
Mapper	Year	Grade
Reducer	Year	AVG(Grade)

	key	value
Mapper	Year	Grade <u>1</u>
Combiner	Year	SUM(Grade) SUM(i)
Reducer	Year	$\frac{\text{SUM(SUMS)}}{\text{SUM(counts)}}$



# Cryptography



# Distributing the Password

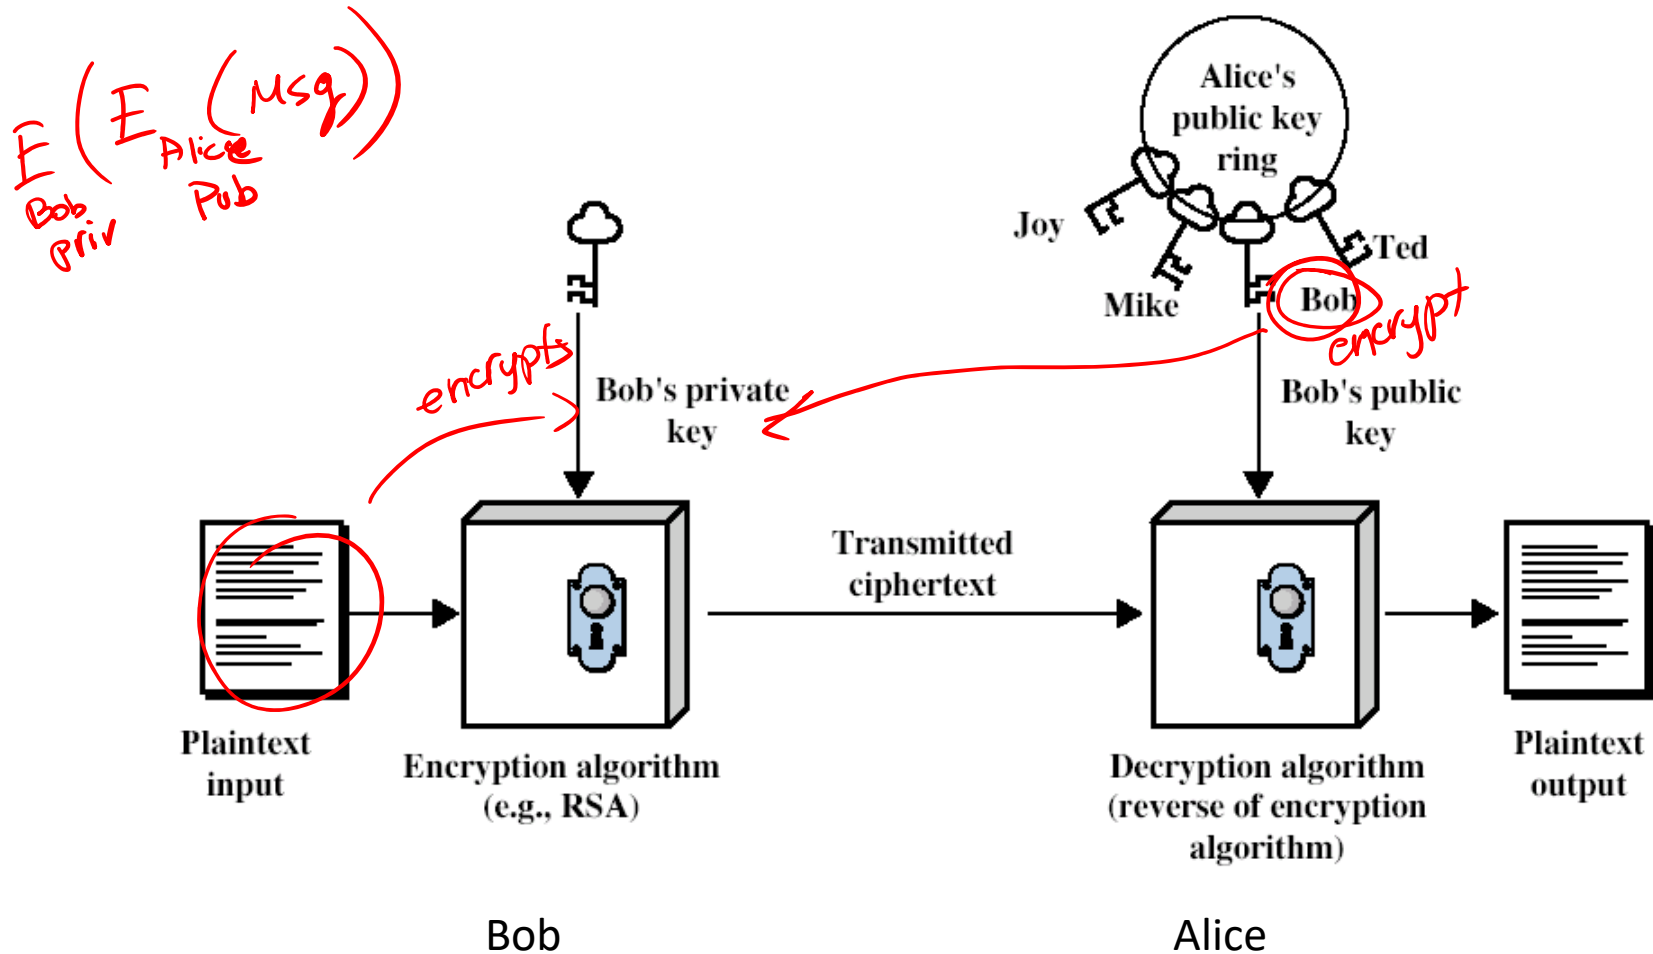
- Password-protected access
- How do you share the password?
  - (without the risk of compromising it)
- Asymmetric encryption!
  - Use a joint pair of keys
  - Public key + Private key

password → Hash(password)

# Public/Private Key

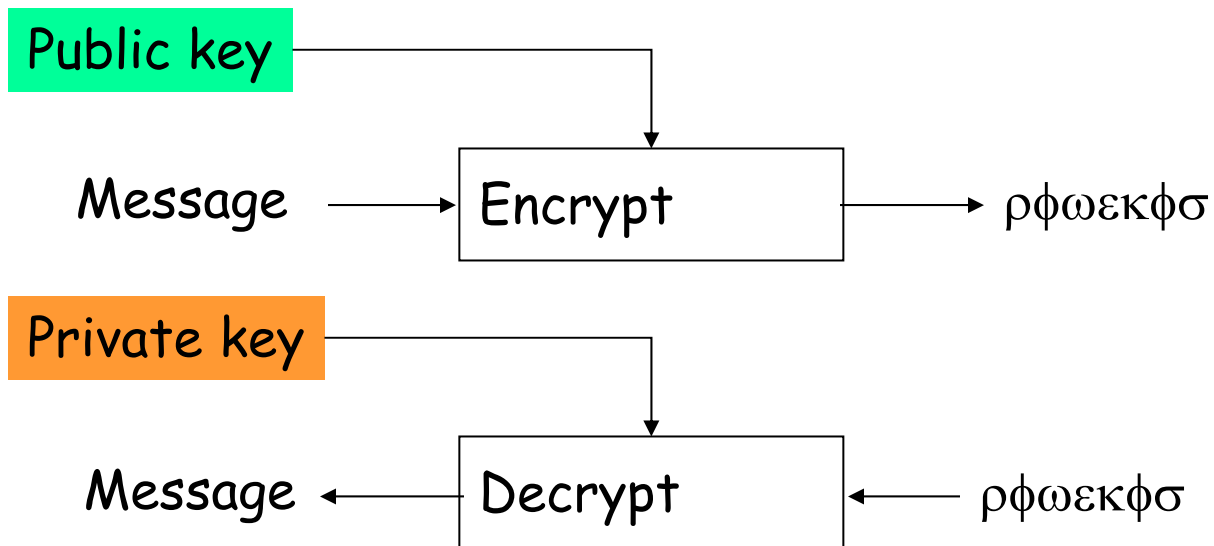
- Private key
  - Known only to the owner
  - Secret code
- Public key
  - Known to everyone and shared freely
    - Need a reliable way to publish
  - Can be used to verify private key authenticity

# Asymmetric Encryption for Authentication



# Public key Encryption

- Alice has a key pair: **public** and **private**
  - **publish** the public key such that the key is publicly known
  - Alice keeps the private key secret
- Other people use Alice's public key to encrypt messages for Alice
- Alice uses her private key to decrypt
- Only Alice can decrypt since only Alice has the private key



- Trick: To compute the private key from the public key is a difficult problem.

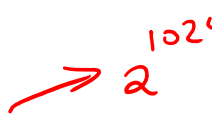




7:40



# RSA

- Invented by Rivest, Shamir & Adleman of MIT in 1977
- Best known and widely used public-key scheme
- Use large integers (e.g. 1024 bits) 
- Security due to cost of factoring large numbers
  - factorization is expensive

# RSA Key Setup

- Each user generates a public/private key pair by
  - select two large primes at random:  $p$ ,  $q$
  - compute their system modulus  $n = p \cdot q$ 
    - note  $\phi(n) = (p-1)(q-1)$
  - select at random the encryption key  $e$ 
    - where  $1 < e < \phi(n)$ ,  $\gcd(e, \phi(n)) = 1$
  - solve following equation to find decryption key  $d$ 
    - $e \cdot d = 1 \pmod{\phi(n)}$  and  $0 \leq d \leq n$
  - publish their public encryption key:  $KU = \{e, n\}$
  - keep secret private decryption key:  $KR = \{d, n\}$

# RSA Usage

- To encrypt a message  $M$ :
  - sender obtains public key of receiver  $KU = \{e, n\}$
  - computes:  $C = M^e \bmod n$ , where  $0 \leq M < n$
- To decrypt the ciphertext  $C$ :
  - receiver uses its private key  $KR = \{d, n\}$
  - computes:  $M = C^d \bmod n$
- Message  $M$  must be smaller than the modulus  $n$  (cut into blocks if needed)

$$M^{ed} = M^{de} = 1 \bmod n$$

# RSA Example: Computing Keys

1. Select primes:  $p=17$ ,  $q=11$
2. Compute  $n=pq=17 \times 11=187$
3. Compute  $\phi(n)=(p-1)(q-1)=16 \times 10=160$
4. Select  $e$ :  $\gcd(e, 160)=1$  and  $e < 160$ 
  - choose  $e=7$
5. Determine  $d$ :  $de=1 \pmod{160}$  and  $d < 160$ 
  - $d=23$  since  $23 \times 7=161=10 \times 160+1$
6. Publish public key  $KU=\{7, 187\}$
7. Keep secret private key  $KR=\{23, 187\}$

# RSA: Encryption and Decryption

- Given message  $M = \underline{88}$  ( $88 < 187$ )

- Encryption  $KU = \{7, 187\}$  :

$$C = \underline{88^7} \bmod 187 = \underline{\underline{11}}$$

*Encrypt  
Bob-public*

- Decryption  $KR = \{23, 187\}$  :

$$M = 11^{\underline{23}} \bmod 187 = \underline{\underline{88}}$$

- $e * d \equiv 1 \pmod{p-1}$  and
- $e * d \equiv 1 \pmod{q-1}$
- Therefore  $m^{ed} \equiv 1 \pmod{pq}$

# Matrix Multiplication

- Matrix times Vector

$$\mathbf{AB} = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} a \times x + b \times y \\ c \times x + d \times y \end{pmatrix} = \begin{pmatrix} ax + by \\ cx + dy \end{pmatrix}$$

- Unit of operation
  - Value-by-value
  - Vector-by-vector

# Example Queries:

```
SELECT StoreID, MIN(TransactionAmt), MAX(TransactionAmt)
FROM StoreTransactions
GROUP BY StoreID;
```

```
SELECT ZipCode, MEDIAN(Age)
FROM People
GROUP BY ZipCode
```

```
SELECT d.category, COUNT(*)
FROM Employee e, Department d
WHERE e.dept = d.id
GROUP BY d.category;
```

*Handwritten red annotations:*  
- "Key #1" next to d.id  
- "Key #2" next to d.category

	key	value
Mapper	Trans ID	TransAmt TransAmt
Combiner	TransID	min(TransAmt) max(TransAmt)
Reducer	TransID	min(TransAmt) max(TransAmt)

Mapper	ZipCode	Age
Reducer	ZipCode	Median(Age)

Mapper 1 E	Dept	- E
Mapper 1 D	ID	Category - D
Reducer k	ID/Dept	Category

Mapper 2	Category	1
Reducer k 2	Category	Sum(1s)



# Hive Execution Flow

- Parse the query
- Get metadata from MetaStore
- Create a logical plan
- Optimize the plan
- Create a physical plan
  - DAG of MR jobs

# Hive Types

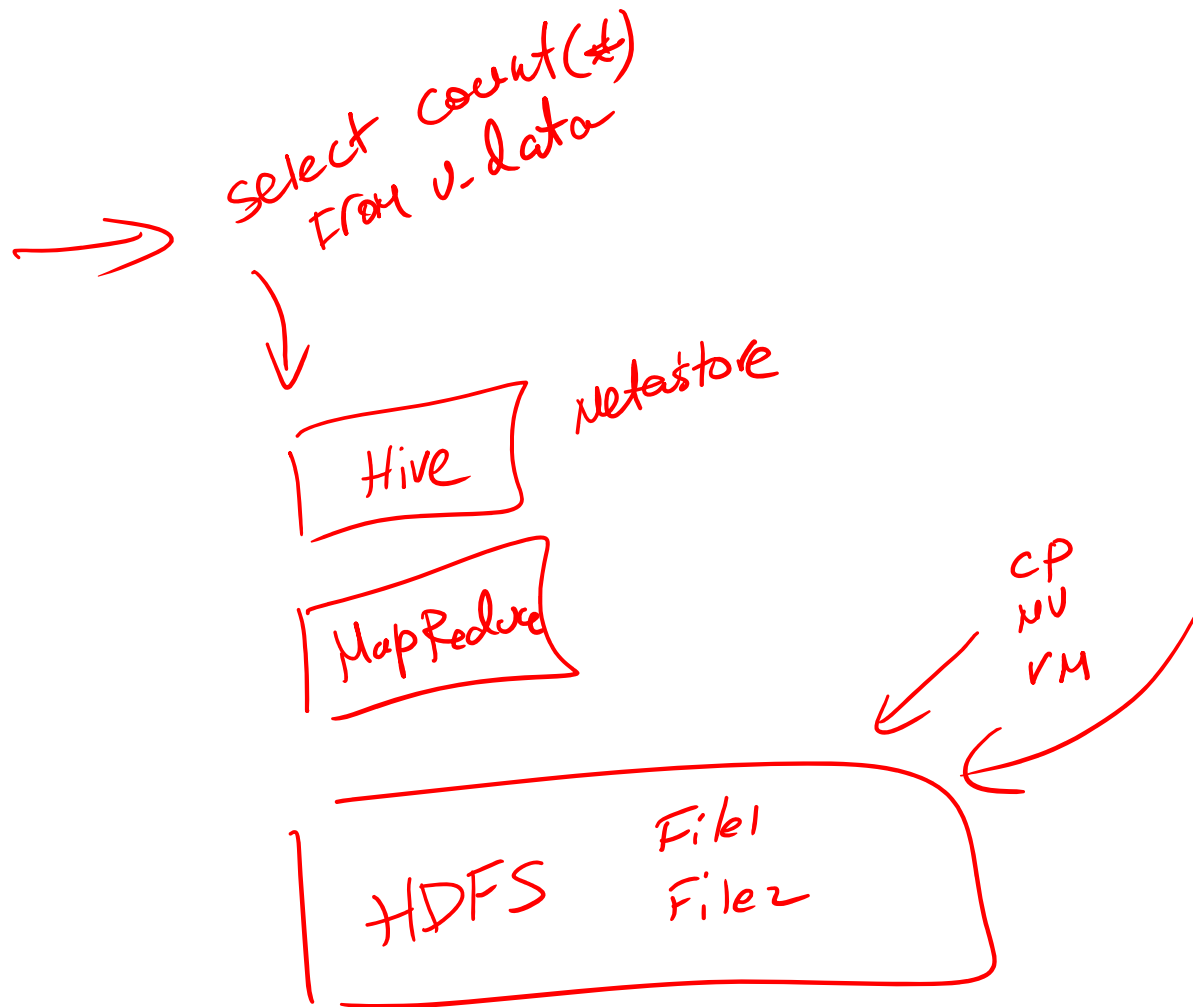
- TINYINT – 1 byte integer
- SMALLINT – 2 byte integer
- INT – 4 byte integer
- BIGINT – 8 byte integer
- BOOLEAN – True/False
- FLOAT/DOUBLE
- STRING

# Hive Examples

```
CREATE TABLE u_data ( userid INT,  
movieid INT,  
rating INT,  
unixtime STRING)  
ROW FORMAT DELIMITED FIELDS  
TERMINATED BY '\t' STORED AS TEXTFILE; (not compressed)
```

- show tables; describe u\_data;
- wget <http://www.grouplens.org/system/files/ml-100k.zip>

```
LOAD DATA LOCAL INPATH 'ml-100k/u.data'  
OVERWRITE INTO TABLE u_data;
```



# Using Hive

- `SELECT COUNT(*) FROM u_data;`
- `SELECT * FROM u_data WHERE userid = 449;`
- `SELECT userid, AVG(rating) from u_data  
GROUP BY userid;`
- `SELECT userid, AVG(rating) as AR from u_data  
GROUP BY userid ORDER BY AR;`

# Custom MapReduce Plug

```
CREATE TABLE u_data_new ( userid INT,  
movieid INT, rating INT, weekday String)  
ROW FORMAT DELIMITED FIELDS  
TERMINATED BY '\t';  
add FILE weekday_mapper.py;  
INSERT OVERWRITE TABLE u_data_new  
SELECT TRANSFORM (userid, movieid, rating, unixtime)  
USING 'python weekday_mapper.py'  
AS (userid, movieid, rating, weekday) FROM u_data;
```

# Using Hive

```
SELECT weekday, COUNT(*)  
FROM u_data_new GROUP BY weekday;
```

```
SELECT weekday, COUNT(*) as Total  
FROM u_data_new GROUP BY weekday  
ORDER BY Total;
```

# Sampling in Hive

```
SELECT COUNT(*) FROM u_data  
TABLESAMPLE(BUCKET 1 OUT OF 100 ON rand());  
  
SELECT COUNT(*) FROM movie_ratings2  
TABLESAMPLE(BUCKET 4 OUT OF 50 ON movieid);  
  
CREATE VIEW MovieSample AS  
SELECT * FROM movie_ratings2  
TABLESAMPLE(0.1 PERCENT);
```



# Joins in Hive

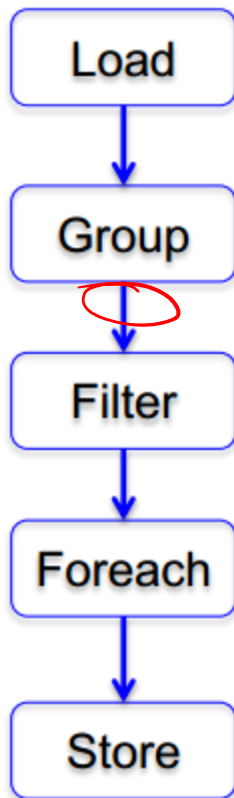
```
SELECT COUNT(*)  
FROM u_data JOIN MovieSample ON  
(u_data.movieid = MovieSample.movieid);  
  
CREATE VIEW JoinView1 AS  
SELECT *  
FROM u_data JOIN MovieSample ON  
(u_data.movieid = u_data_new.movieid)  
WHERE Rating > 3;
```

# Outer Join / External Table

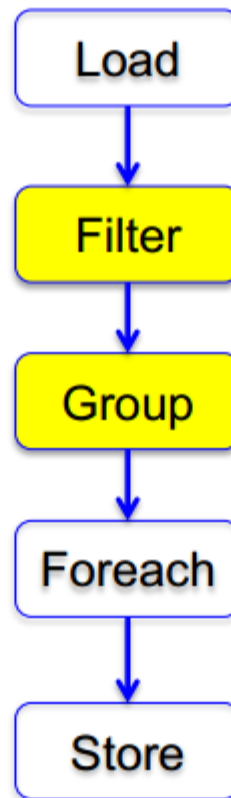
- External Tables
  - Data not owned by Hive
- Describe u\_data;
- Describe u\_data\_new;
- ```
SELECT * FROM u_data
FULL OUTER JOIN u_data_new ON
(u_data.rating = u_data_new.weekday)
WHERE u_data.rating > 3;
```

# Pig Architecture

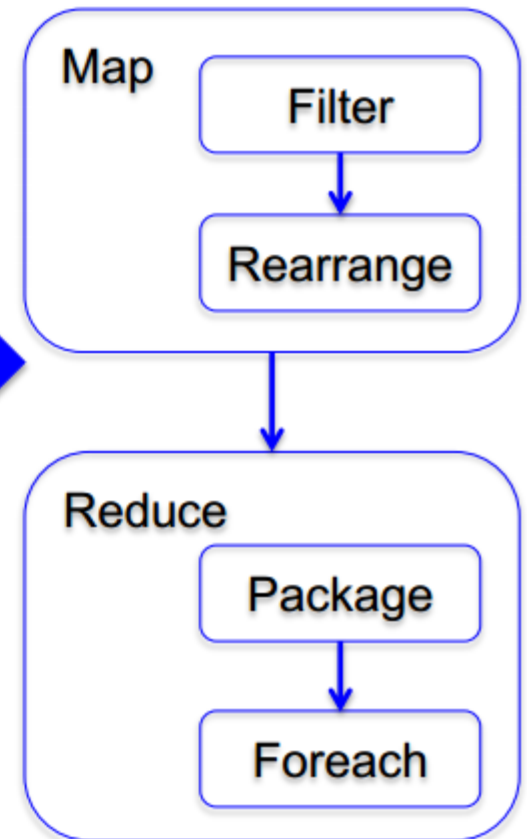
Logical Plan



Rule  
based  
optimizations



MapReduce Plan



# Pig Use

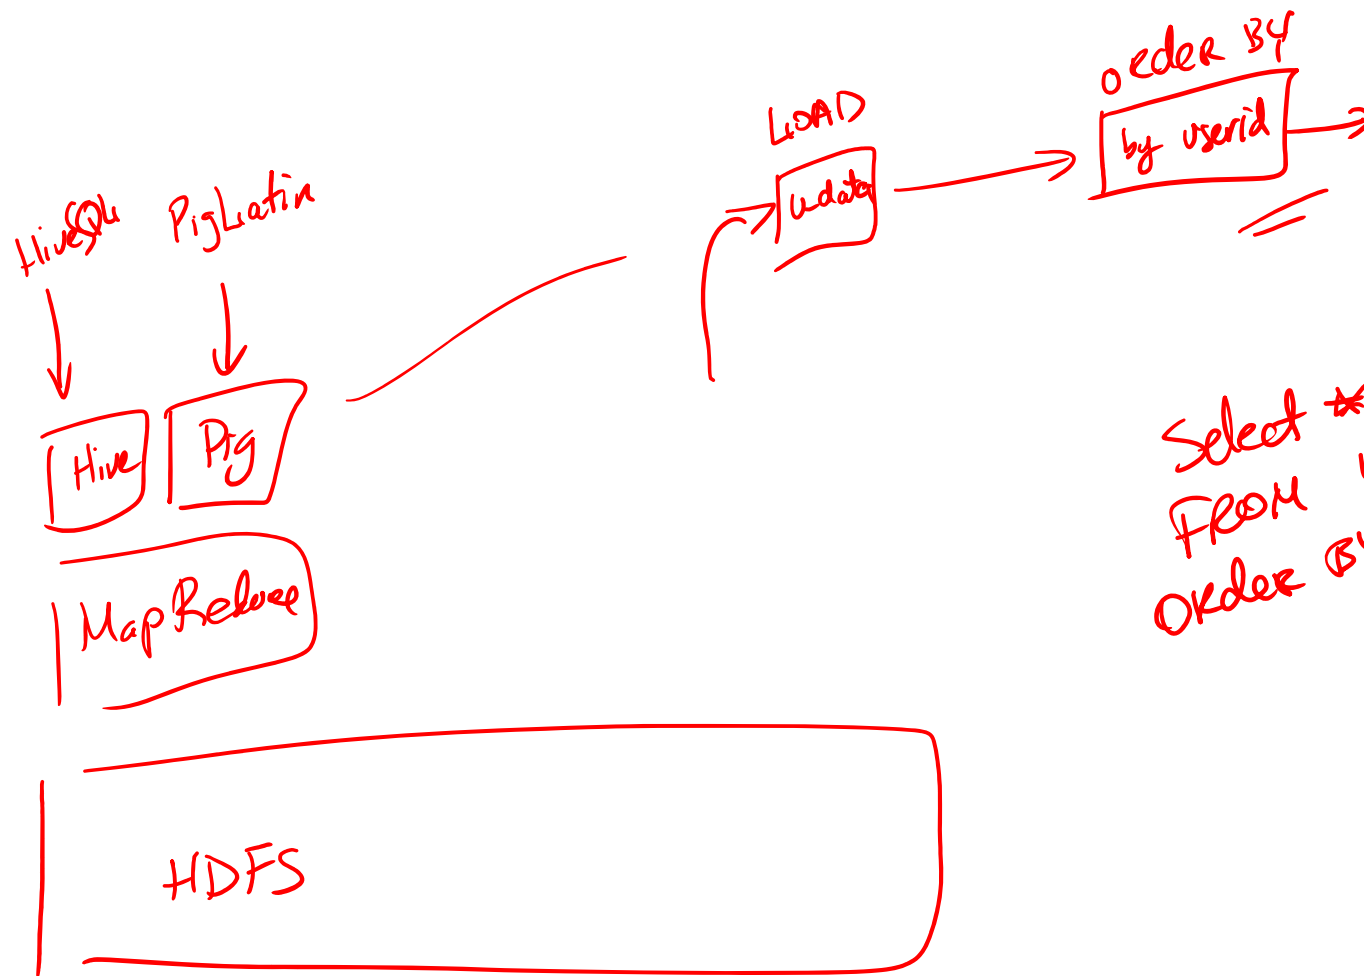
```
UData = LOAD 'u.data' USING PigStorage('\t') AS  
(userid:int, movieid:int, rating:int,  
unixtime:chararray);
```

```
DESCRIBE UData;
```

```
DUMP UData;
```

```
UDataS = ORDER UData BY userid;
```

```
DUMP UDataS;
```



Select \*  
FROM v-data  
Order BY userid

# We Have Stopped Here

# Pig Use

```
UDataSample = SAMPLE UData 0.01;  
DUMP UDataSample;  
STORE UDataSample INTO 'UDataSample'  
USING PigStorage ('_');
```

```
ONE_USER = FILTER UData BY userid == 251;  
DUMP ONE_USER;
```

# Pig Use

```
GoodRatings = FILTER UData BY rating > 2;  
UserSet = GROUP GoodRatings BY userid;  
DUMP UserSet;  
UserRatings = FOREACH UserSet GENERATE  
COUNT(GoodRatings);  
DUMP UserRatings;  
ILLUSTRATE UserRatings;
```



# Pig Use

```
UserSet2 = GROUP UData BY userid;  
UserRatings2 = FOREACH UserSet2 GENERATE  
UData.userid, AVG(UData.rating);  
DUMP UserRatings2;
```

# Next Time:

- Hadoop ecosystem
  - Hadoop config, Hadoop Streaming
  - More Hive and Pig
- Read:
  - Mining of Massive Datasets
    - Sections 2.5
  - Hadoop: The Definitive Guide
    - Pp162-167, "Launching a Job" to "Retrieving the Results"
    - Pp204-205, "Speculative Execution"