

CSC 555 and DSC 333

Mining Big Data

Lecture 6

Alexander Rasin

College of CDM, DePaul University

October 19th, 2021

Tonight

- Performance considerations and compression
- NoSQL
- Multi-node cluster setup
- Hadoop streaming join

Performance

- Query cost
 - Load data *put data into HDFS*
 - Map-time
 - Reduce-time
 - Tuning actions
 - ✓ – Replication *1x → 3x*
 - Compression
 - Balancing: blocks and keys
-

Replication

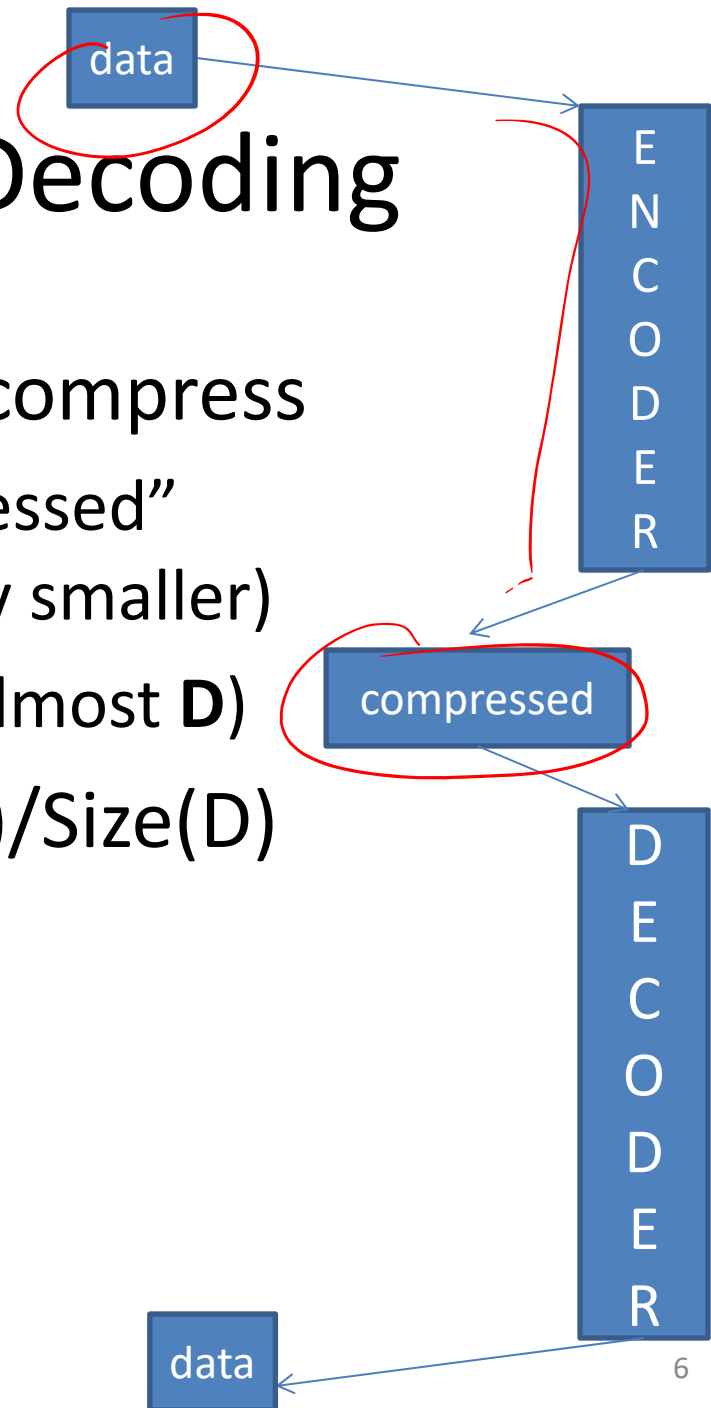
- Number of times a block is copied in HDFS
- Expected performance change?
 - Load
 - Map-Runtime
 - Reduce-Runtime
- Configuration (JobConf)
 - ✓ – setMapSpeculativeExecution
 - setReduceSpeculativeExecution

Data Compression

- Reduce the data set (file)
 - Save space when storing the file
 - Save time when reading/sending the file
 - Pay overhead for compression/decompression
- Applications
 - Gzip, Zip, etc.
 - MP3, JPEG, MPEG, ...
 - Built-in (databases, file systems, etc.)

Encoding and Decoding

- Data: binary data set **D** to compress
 - Encode: generate a “compressed” representation D_c (hopefully smaller)
 - Decode: reconstruct **D** (or almost **D**)
- Compression ratio: $\text{Size}(D_c)/\text{Size}(D)$
- Lossless/Lossy
 - Text, files
 - Images, MP3, Video, ...




Simple Ideas

- Run Length Encoding (RLE)
 - Convert (AAAAACCCCCDDDDDDDD)
 - Into (A,5;C,4;D,7)
- Fixed length (dictionary) coding
 - Assign codes to every letter
 - E.g., a = 00001, b = 00010, ... z = 11000
 - Encode character => code
 - Have to store the mapping

Variable Length Encoding

- Fewer bits for frequent letters (E, T, A)
- More bits for rare letters (J, Q, Z)
- Potential ambiguity
 - (E=0, T=01, Z=0001101)?
- Use a binary tree



char	encoding
a	1
b	001
c	0000
d	0001
r	01

- *a a b c r r d a*
 110010000010100011

Compression Quality


- Huffman coding/Huffman tree
 - Build a tree based on character frequency
- Difficult to compress
 - Random data (or compressed data)
- All compression algorithms exploit bias in data
 - Images have white patches
 - Letters E/T/A occur more frequently

Entropy. (Shannon 1948)
$$H(S) = \sum_{s \in S} p(s) \log_2 \frac{1}{p(s)}$$

2021 Hadoop 1
2021 Hadoop 1 → 2021 Hadoop 2


Compression in Hadoop

- Multiple (independent) opportunities
 - HDFS => Mapper => Reducer => HDFS



Assigning Data to Mappers

- HDFS-based
- Mapper executed locally for best performance
- Everything is managed in blocks
 - Replication
 - Balancing



Block Rebalancing

- How the data is distributed on HDFS nodes
- Expected performance change?
 - Load
 - Map-Runtime
 - Reduce-Runtime

Running the Rebalancer

- Runs in background
- Specify the threshold
 - `hdfs balancer -threshold X`
 - Percent of node usage
 - (default 10%)
- Configuration
 - `dfs.balance.bandwidthPerSec` limits the amount of data each node spends on rebalancing (1MB default)

Assigning Data to Reducers

- Partitioner
 - `JobConf.setPartitionerClass()`
- Uses Java hash function + MOD by default
- Sub any function (based on # of reducers)

Partitioning Function

- Function assigning Key-value pairs to Reducers
- Expected performance change?
 - Load
 - Map-Runtime
 - Reduce-Runtime
- Sorting



- No SQL / Not Only SQL
- Non-relational databases (data-store) systems
- Relax some of the relational database rules
 - Integrity constraints
 - Rigid schema (some)
 - Update/synchronization guarantees
 - Flexible
 - Highly-scalable
- Always a trade-off

Key Value Stores

- For a (key, value) combination
 - Key is a key
 - Value is a “blob”
 - Insert
 - Update
 - Delete
 - Get
- Many implementations
 - DynamoDB, Riak, Redis, BerkeleyDB, ...

Document Stores

{ "Author": "Rusty" }

- Similar to key value stores
- Value is not a blob, but a document, e.g.
 - { "Subject": "I like Plankton"
"Author": "Rusty"
"PostedDate": "5/23/2006"
"Tags": ["plankton", "baseball", "decisions"]
"Body": "I decided today that I don't like baseball. I like plankton." }
- You may have heard of
 - MongoDB, CouchDB, Terrastore, RavenDB

Column Family Stores

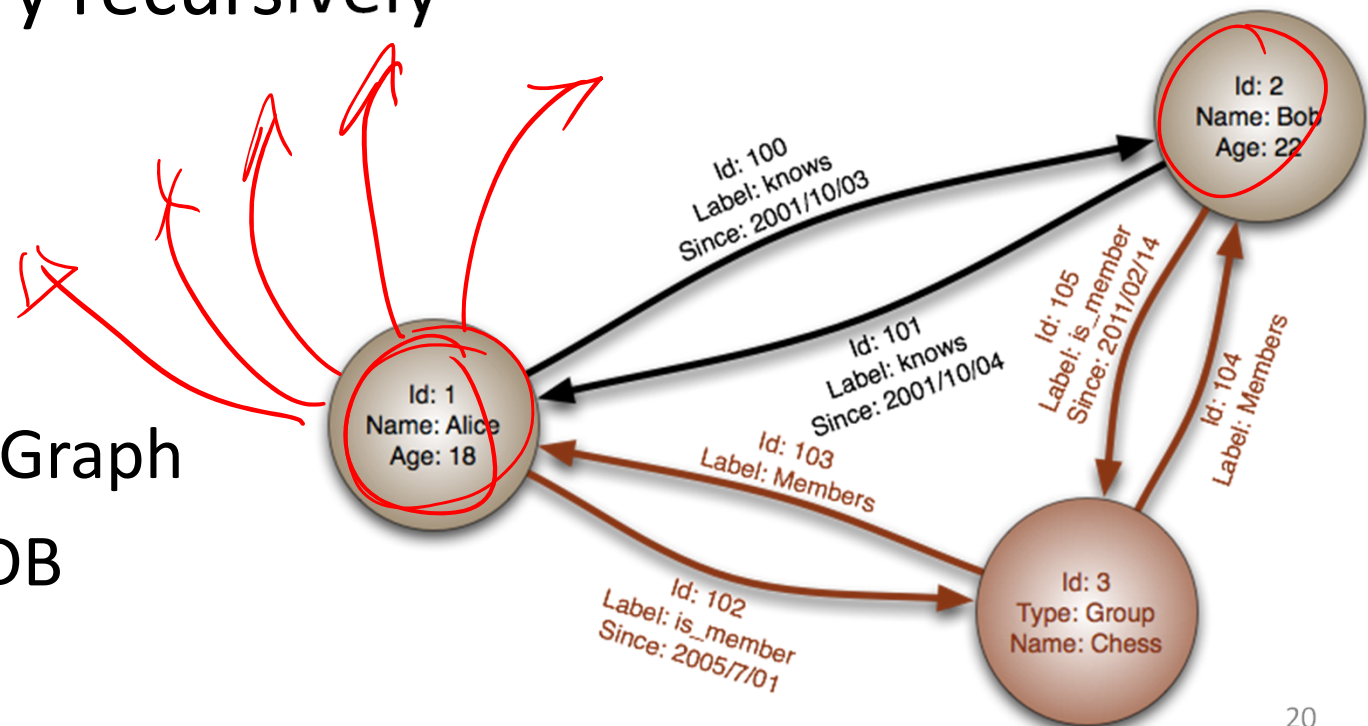
- Similar to document stores
- Add the concept of “column family”
- Helps optimize read access and updates

	F1		F2		F3		
	C1	C2	C5	C6	C7	C8	C9
1	(1)	(2)	(3)	(4)			
2							
3							
4							
5							
6							
7							

Graph Databases

- Database stores the set of items
- Items are connected to each other
- Can query recursively

- DBs
 - Neo4j
 - InfiniteGraph
 - OrientDB



7:45


A Break




Host Names/IP

- Nodes addressed by
 - Internet Protocol (IP) Address
 - 140.192.5.61 *IPv4*
 - Domain Name
 - www.depaul.edu
- IP and Domain name are interchangeable


Domain Reverse Hierarchy



DNS database




UNIX filesystem



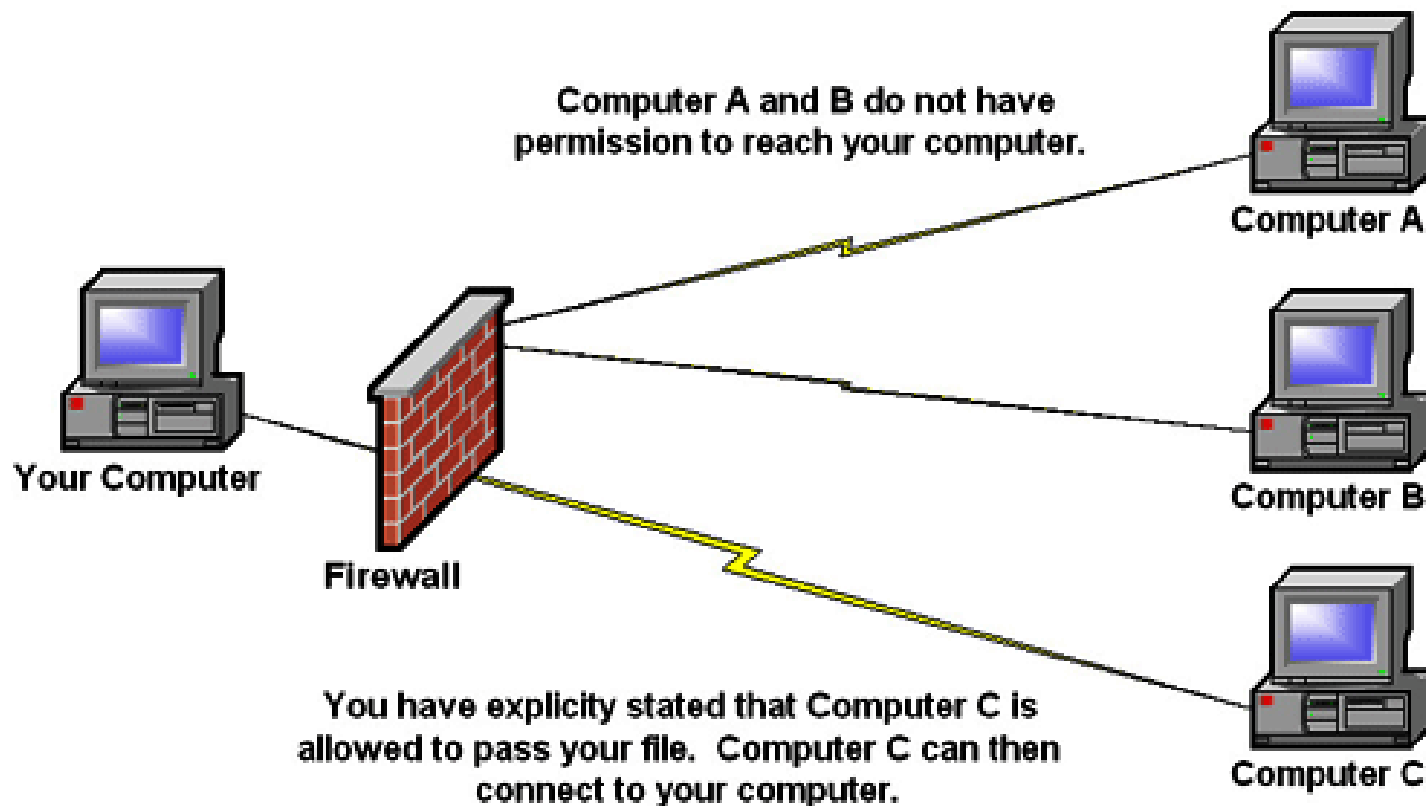
Machine Address + Port

- Address always includes a port
 - Sometimes not explicitly specified
 - Well known defaults
 - ssh = 22
 - www = 80
 - ftp = 21
 - mail = 25
 - Lesser known
 - Hadoop status report = 50070




Firewall

- Controls the individual port access



Cluster Setup

- Passwordless login
 - Public/private key
 - Localhost to Localhost
 - Master to Worker
- Firewall / Security groups
- Hadoop logs



Hadoop Streaming: A Join!

```
SELECT SUBSTRING(tv.Plate, 0, 3), COUNT(*)  
FROM TowedVehicles AS tv,  
      RelocatedVehicles AS rv  
WHERE  
SUBSTRING(tv.Plate, 0, 3) = SUBSTRING(rv.Plate, 0, 3)  
GROUP BY SUBSTRING(tv.Plate, 0, 3)
```

Hive Transform

```
CREATE TABLE u_data ( userid INT, movieid INT, rating INT,  
unixtime STRING) ROW FORMAT DELIMITED FIELDS  
TERMINATED BY '\t' STORED AS TEXTFILE; (not compressed)
```

```
LOAD DATA LOCAL INPATH 'ml-100k/u.data'  
OVERWRITE INTO TABLE u_data;
```

```
INSERT OVERWRITE TABLE u_data_new  
SELECT TRANSFORM (userid, movieid, rating, unixtime) USING  
'python weekday_mapper.py'  
AS (userid, movieid, rating, weekday) FROM u_data;
```

Next Time:

- More on Hadoop Ecosystem
- Mahout
- Link analysis